

# **COMP9313: Big Data Management**



**Lecturer: Xin Cao**

**Course web site:** <http://www.cse.unsw.edu.au/~cs9313/>

# **Chapter 5: Graph Data Processing in MapReduce**

# What's a Graph?

- $G = (V, E)$ , where
  - $V$  represents the set of vertices (nodes)
  - $E$  represents the set of edges (links)
  - Both vertices and edges may contain additional information
- Different types of graphs:
  - Directed vs. undirected edges
  - Presence or absence of cycles
- Graphs are everywhere:
  - Hyperlink structure of the Web
  - Physical structure of computers on the Internet
  - Interstate highway system
  - Social networks

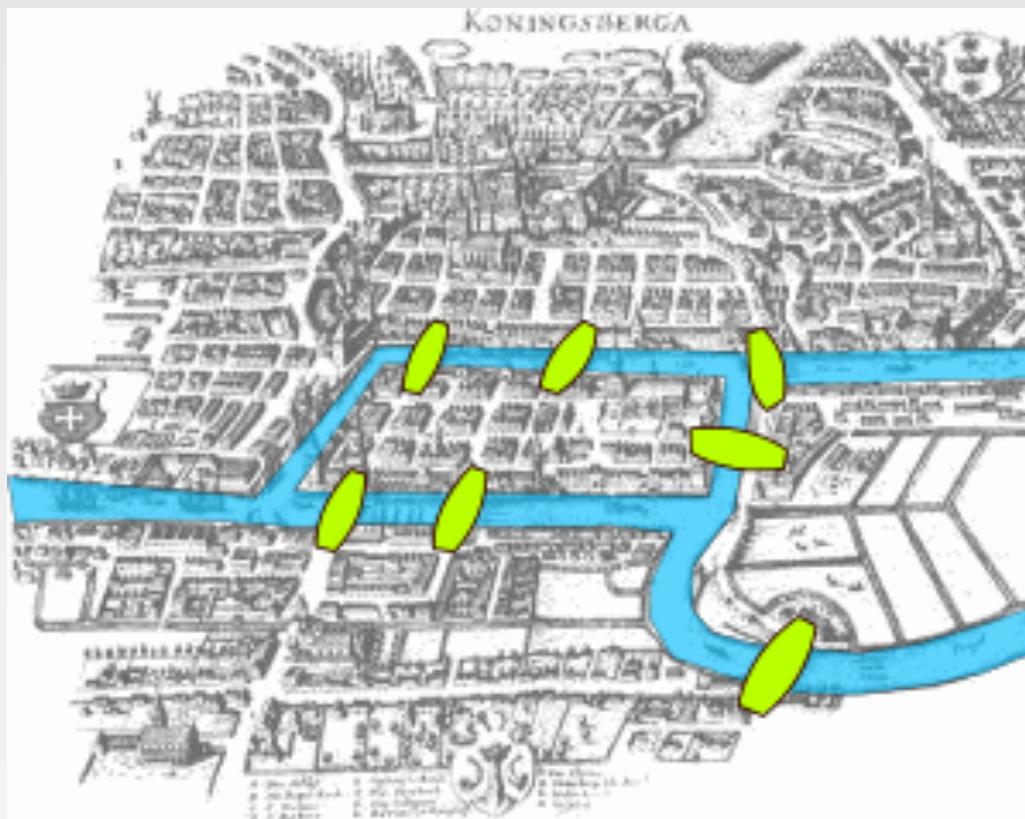
# Graph Data: Social Networks



Facebook social graph

4-degrees of separation [Backstrom-Boldi-Rosa-Ugander-Vigna, 2011]

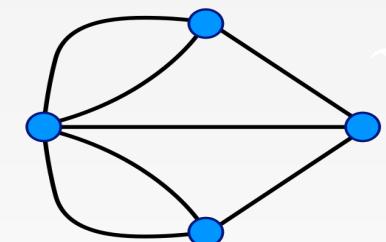
# Graph Data: Technological Networks



## Seven Bridges of Königsberg

[Euler, 1735]

Return to the starting point by traveling each link of the graph once and only once.



# Some Graph Problems

- Finding shortest paths
  - Routing Internet traffic and UPS trucks
- Finding minimum spanning trees
  - Telco laying down fiber
- Finding Max Flow
  - Airline scheduling
- Identify “special” nodes and communities
  - Breaking up terrorist cells, spread of avian flu
- Bipartite matching
  - Monster.com, Match.com
- And of course... **PageRank**

# Graph Analytics

- General Graph
  - Count the number of nodes whose degree is equal to 5
  - Find the diameter of the graphs
- Web Graph
  - Rank each webpage in the web graph or each user in the twitter graph using PageRank, or other centrality measure
- Transportation Network
  - Return the shortest or cheapest flight/road from one city to another
- Social Network
  - Detect a group of users who have similar interests
- Financial Network
  - Find the path connecting two suspicious transactions;
- ....

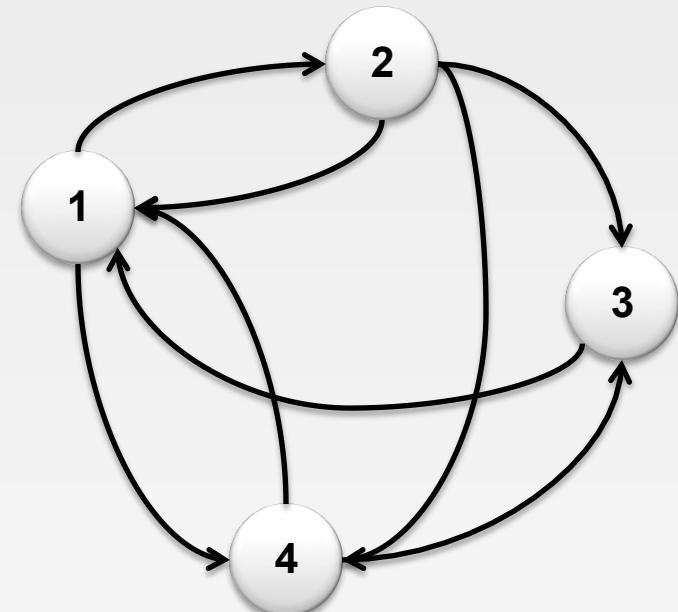
# Graphs and MapReduce

- Graph algorithms typically involve:
  - Performing computations at each node: based on node features, edge features, and local link structure
  - Propagating computations: “traversing” the graph
- Key questions:
  - How do you represent graph data in MapReduce?
  - How do you traverse a graph in MapReduce?

# Representing Graphs

- Adjacency Matrices: Represent a graph as an  $n \times n$  square matrix  $M$ 
  - $n = |V|$
  - $M_{ij} = 1$  means a link from node  $i$  to  $j$

	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	1	0	0	0
4	1	0	1	0



# Adjacency Matrices: Critique

## ■ Advantages:

- Amenable to mathematical manipulation
- Iteration over rows and columns corresponds to computations on outlinks and inlinks

## ■ Disadvantages:

- Lots of zeros for sparse matrices
- Lots of wasted space

# Representing Graphs

- Adjacency Lists: Take adjacency matrices... and throw away all the zeros

	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	1	0	0	0
4	1	0	1	0



1: 2, 4  
2: 1, 3, 4  
3: 1  
4: 1, 3

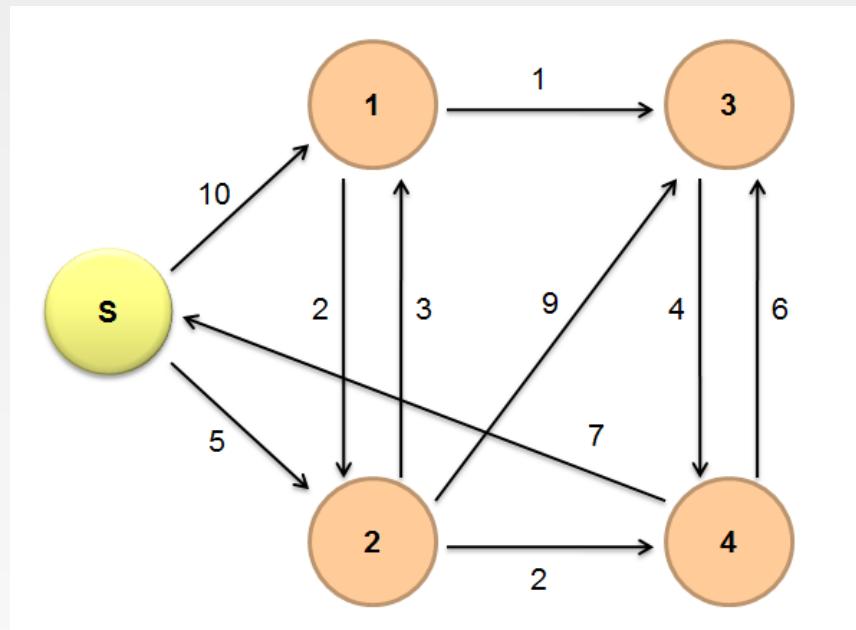
# Adjacency Lists: Critique

- Advantages:
  - Much more compact representation
  - Easy to compute over outlinks
- Disadvantages:
  - Much more difficult to compute over inlinks

# **Single-Source Shortest Path**

# Single-Source Shortest Path (SSSP)

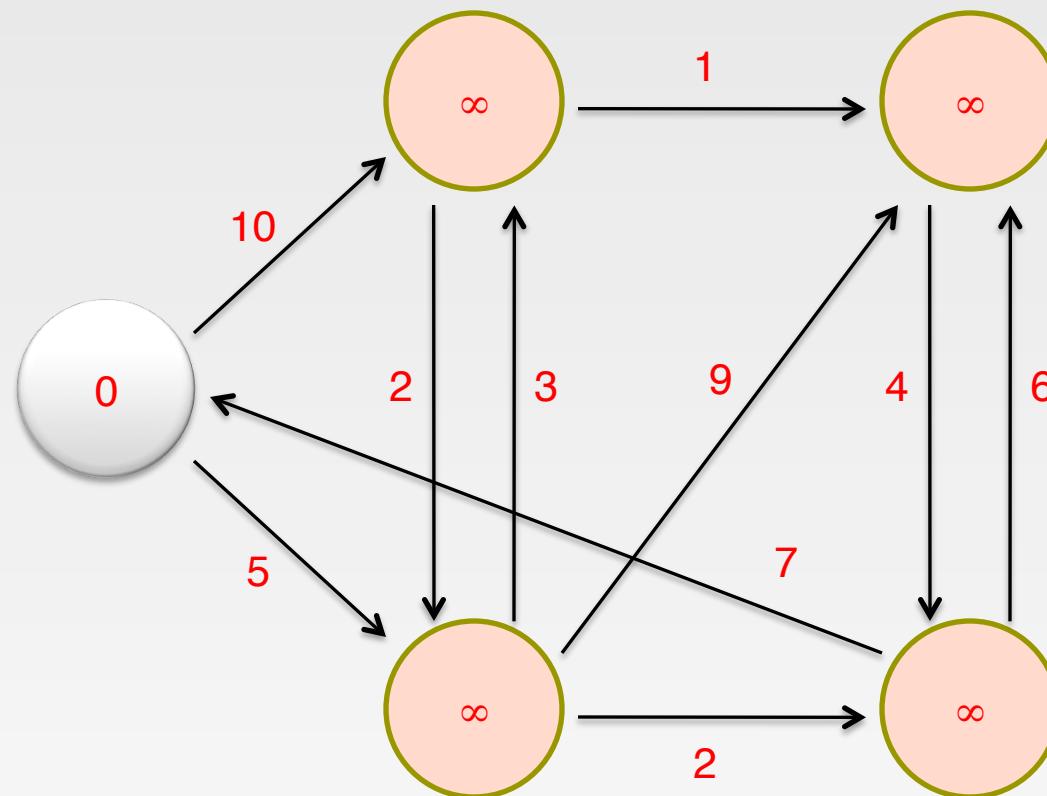
- **Problem:** find shortest path from a source node to one or more target nodes
  - Shortest might also mean lowest weight or cost
- Dijkstra's Algorithm:
  - For a given source node in the graph, the algorithm finds the shortest path between that node and every other



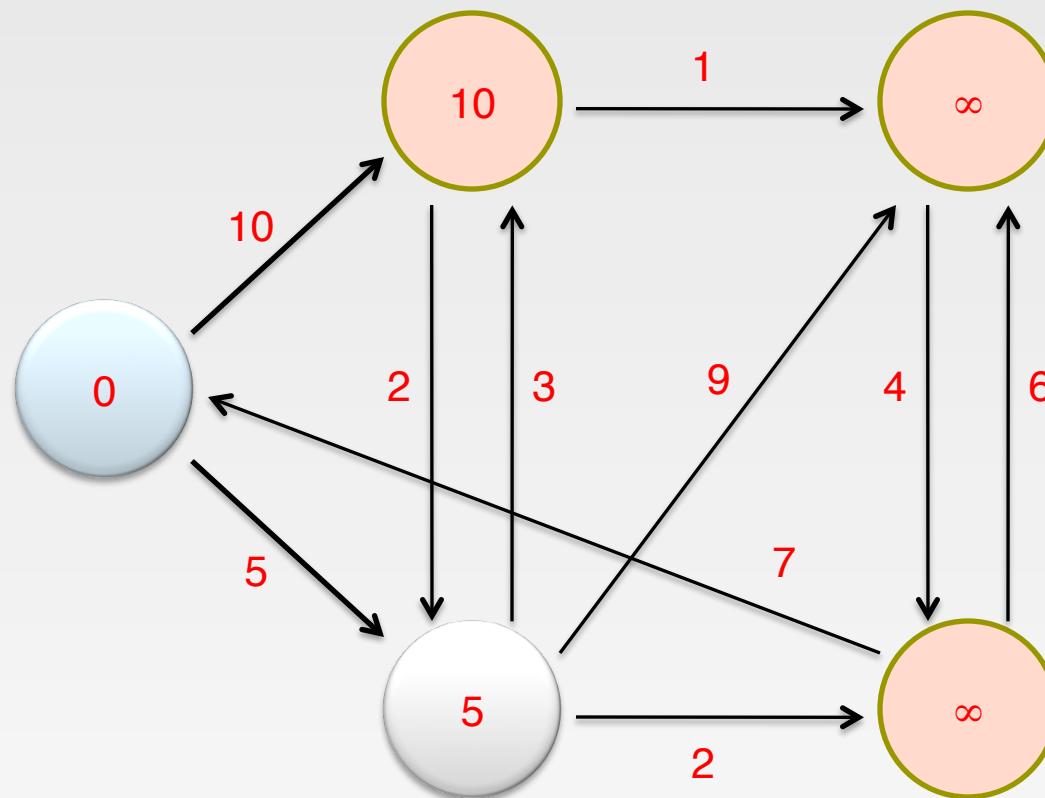
# Dijkstra's Algorithm

```
1: DIJKSTRA( $G, w, s$ )
2:    $d[s] \leftarrow 0$ 
3:   for all vertex  $v \in V$  do
4:      $d[v] \leftarrow \infty$ 
5:    $Q \leftarrow \{V\}$ 
6:   while  $Q \neq \emptyset$  do
7:      $u \leftarrow \text{EXTRACTMIN}(Q)$ 
8:     for all vertex  $v \in u.\text{ADJACENCYLIST}$  do
9:       if  $d[v] > d[u] + w(u, v)$  then
10:         $d[v] \leftarrow d[u] + w(u, v)$ 
```

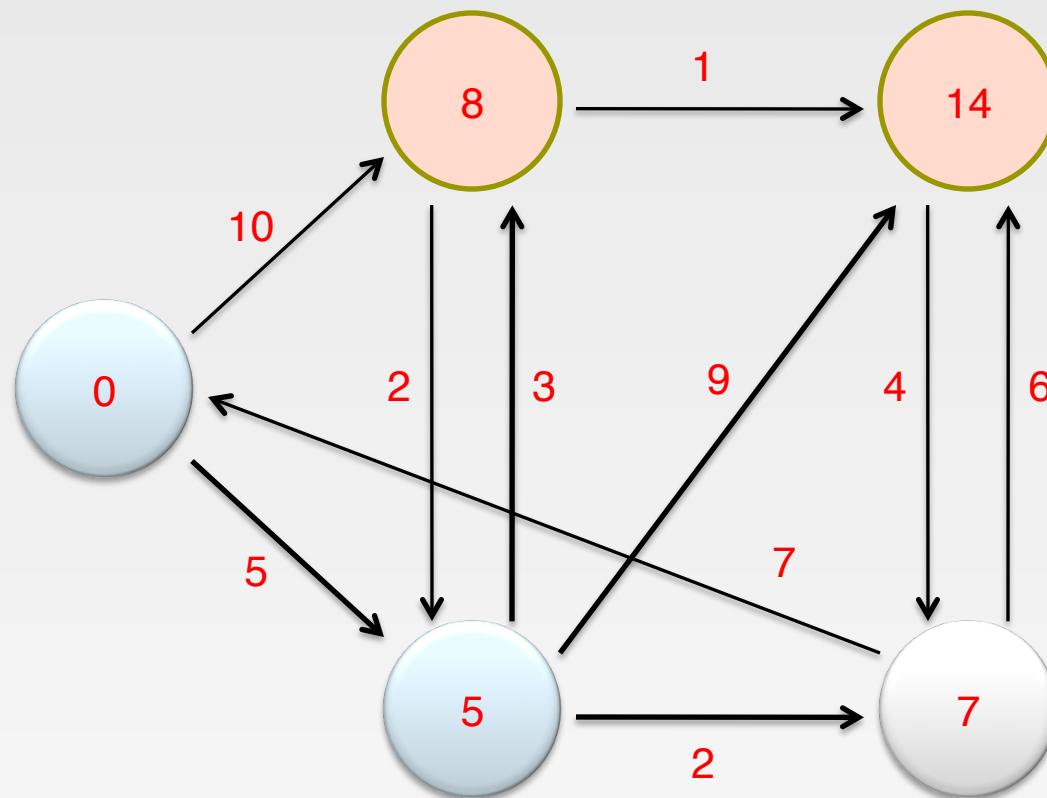
# Dijkstra's Algorithm Example



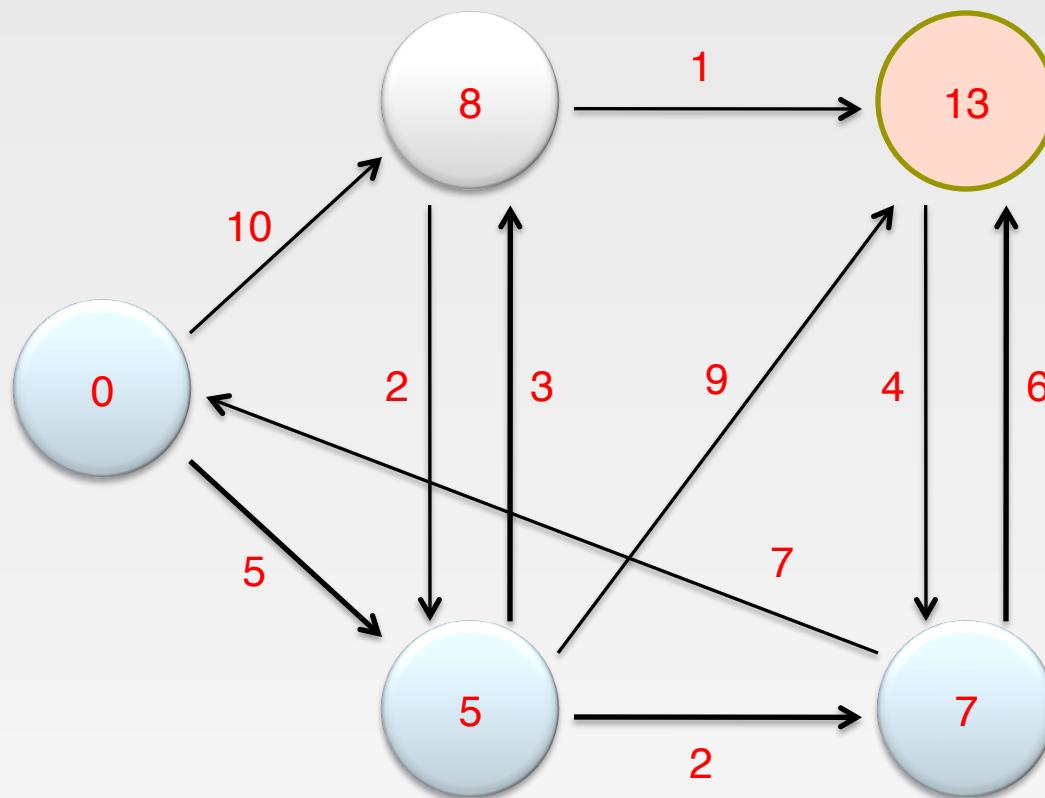
# Dijkstra's Algorithm Example



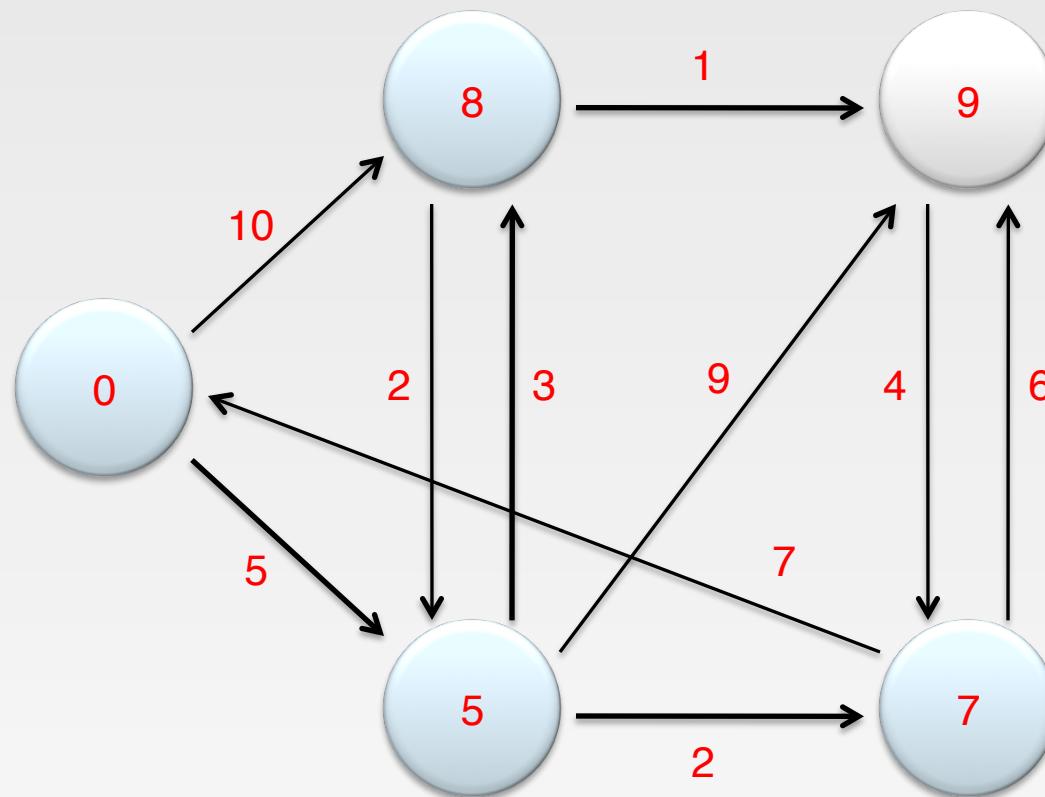
# Dijkstra's Algorithm Example



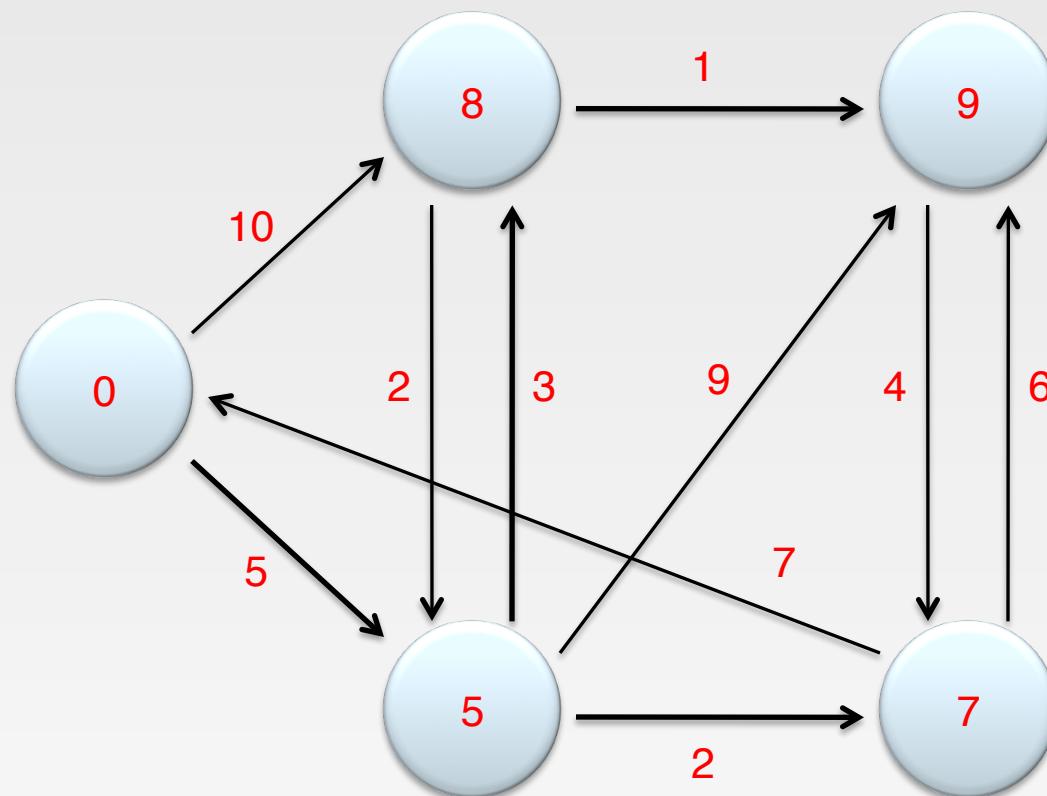
# Dijkstra's Algorithm Example



# Dijkstra's Algorithm Example



# Dijkstra's Algorithm Example



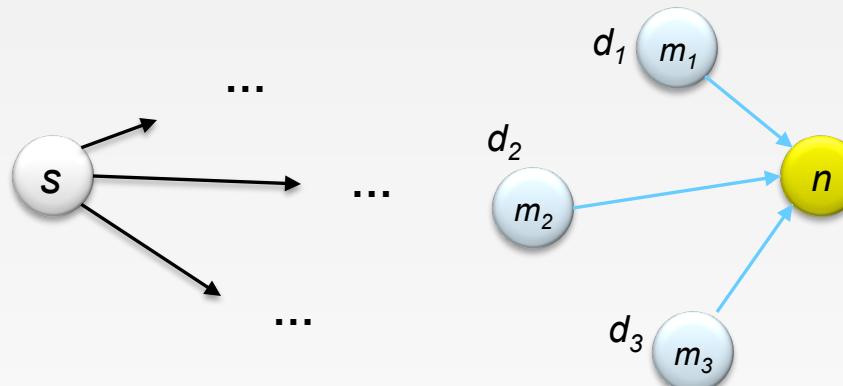
Finish!

# Single Source Shortest Path

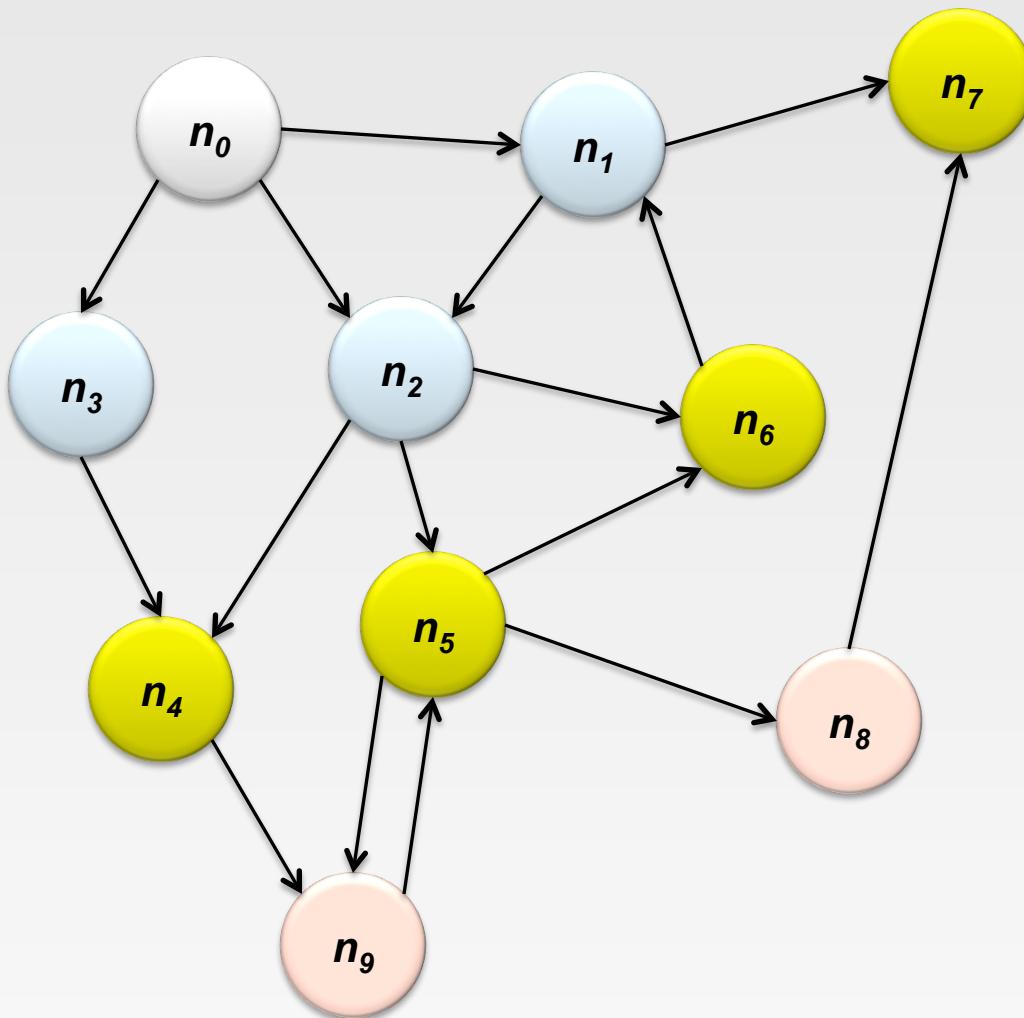
- **Problem:** find shortest path from a source node to one or more target nodes
  - Shortest might also mean lowest weight or cost
- Single processor machine: Dijkstra's Algorithm
- MapReduce: parallel Breadth-First Search (BFS)

# Finding the Shortest Path

- Consider simple case of equal edge weights
- Solution to the problem can be defined inductively
- Here's the intuition:
  - Define:  $b$  is reachable from  $a$  if  $b$  is on adjacency list of  $a$
  - ★  $\text{DISTANCETO}(s) = 0$
  - For all nodes  $p$  reachable from  $s$ ,  
 $\text{DISTANCETO}(p) = 1$
  - For all nodes  $n$  reachable from some other set of nodes  $M$ ,  
 $\text{DISTANCETO}(n) = 1 + \min(\text{DISTANCETO}(m), m \in M)$



# Visualizing Parallel BFS



# From Intuition to Algorithm

- Data representation:
  - Key: node  $n$
  - Value:  $d$  (distance from start), adjacency list (list of nodes reachable from  $n$ )
  - Initialization: for all nodes except for start node,  $d = \infty$
- Mapper:
  - $\forall m \in$  adjacency list: emit  $(m, d + 1)$
- Sort/Shuffle
  - Groups distances by reachable nodes
- Reducer:
  - Selects minimum distance path for each reachable node
  - Additional bookkeeping needed to keep track of actual path

# Multiple Iterations Needed

- Each MapReduce iteration advances the “known frontier” by one hop
  - Subsequent iterations include more and more reachable nodes as frontier expands
  - The input of Mapper is the output of Reducer in the previous iteration
  - Multiple iterations are needed to explore entire graph
- Preserving graph structure:
  - Problem: Where did the adjacency list go?
  - Solution: mapper emits  $(n, \text{adjacency list})$  as well

# BFS Pseudo-Code

- Equal Edge Weights ([how to deal with weighted edges?](#))
- Only distances, no paths stored ([how to obtain paths?](#))

```
class Mapper
    method Map(nid n, node N)
        d ← N.Distance
        Emit(nid n, N)           //Pass along graph structure
        for all nodeid m ∈ N.AdjacencyList do
            Emit(nid m, d+1)     //Emit distances to reachable nodes
```

```
class Reducer
    method Reduce(nid m, [d1, d2, . . .])
        dmin ← ∞
        M ← ∅
        for all d ∈ counts [d1, d2, . . .] do
            if IsNode(d) then
                M ← d           //Recover graph structure
            else if d < dmin then
                dmin ← d       //Look for shorter distance
            M.Distance ← dmin //Update shortest distance
        Emit(nid m, node M)
```

# Stopping Criterion

- How many iterations are needed in parallel BFS (equal edge weight case)?
- Convince yourself: when a node is first “discovered”, we’ve found the shortest path
- Now answer the question...
  - The diameter of the graph, or the greatest distance between any pair of nodes
  - Six degrees of separation?
    - ▶ If this is indeed true, then parallel breadth-first search on the global social network would take at most six MapReduce iterations.

# Implementation in MapReduce

- The actual checking of the termination condition must occur outside of MapReduce.
- The driver (main) checks to see if a termination condition has been met, and if not, repeats.
- Hadoop provides a lightweight API called “counters”.
  - It can be used for counting events that occur during execution, e.g., number of corrupt records, number of times a certain condition is met, or anything that the programmer desires.
  - Counters can be designed to count the number of nodes that have distances of  $\infty$  at the end of the job, the driver program can access the final counter value and check to see if another iteration is necessary.

# How to Find the Shortest Path?

- The parallel breadth-first search algorithm only finds the shortest distances.
- Store “back-pointers” at each node, as with Dijkstra's algorithm
  - Not efficient to recover the path from the back-pointers
- A simpler approach is to emit paths along with distances in the mapper, so that each node will have its shortest path easily accessible at all times
  - The additional space requirement is acceptable

# Mapper

```
public static class TheMapper extends Mapper<LongWritable, Text, LongWritable, Text> {
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        Text word = new Text();
        String line = value.toString(); //looks like 1 0 2:3:
        String[] sp = line.split(" "); //splits on space
        int distanceadd = Integer.parseInt(sp[1]) + 1;
        String[] PointsTo = sp[2].split(":");
        for(int i=0; i<PointsTo.length; i++){
            word.set("VALUE "+distanceadd); //tells me to look at distance value
            context.write(new LongWritable(Integer.parseInt(PointsTo[i])), word);
            word.clear(); }
        //pass in current node's distance (if it is the lowest distance)
        word.set("VALUE "+sp[1]);
        context.write( new LongWritable( Integer.parseInt( sp[0] ) ), word );
        word.clear();
        word.set("NODES "+sp[2]); //tells me to append on the final tally
        context.write( new LongWritable( Integer.parseInt( sp[0] ) ), word );
        word.clear();
    }
}
```

# Reducer

```
public static class TheReducer extends Reducer<LongWritable, Text, LongWritable, Text> {
    public void reduce(LongWritable key, Iterable<Text> values, Context context) throws
IOException, InterruptedException {
        String nodes = "UNMODED";
        Text word = new Text();
        int lowest = INFINITY;//start at infinity

        for (Text val : values) {
            //looks like NODES/VALUES 1 0 2:3;, we need to use the first as a key
            String[] sp = val.toString().split(" ");//splits on space
            //look at first value
            if(sp[0].equalsIgnoreCase("NODES")){
                nodes = null;
                nodes = sp[1];
            }else if(sp[0].equalsIgnoreCase("VALUE")){
                int distance = Integer.parseInt(sp[1]);
                lowest = Math.min(distance, lowest);
            }
        }
        word.set(lowest+" "+nodes);
        context.write(key, word);
        word.clear();
    }
}
```

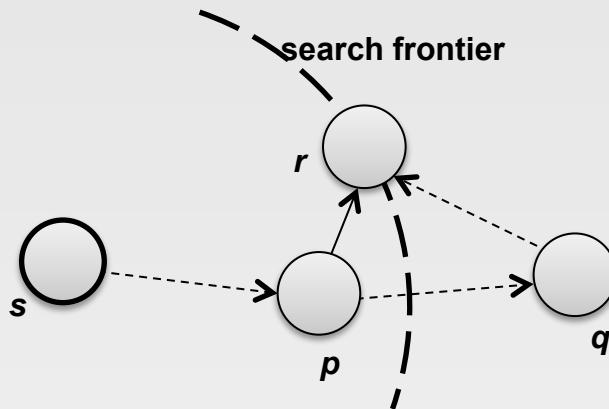
<https://github.com/himank/Graph-Algorithm-MapReduce/blob/master/src/DijikstraAlgo.java>

# BFS Pseudo-Code (Weighted Edges)

- The adjacency lists, which were previously lists of node ids, must now encode the edge distances as well
  - Positive weights!
- In line 6 of the mapper code, instead of emitting  $d + 1$  as the value, we must now emit  $d + w$ , where  $w$  is the edge distance
- **The termination behaviour is very different!**
  - How many iterations are needed in parallel BFS (positive edge weight case)?
  - Convince yourself: when a node is first “discovered”, we’ve found the shortest path

*Not true!*

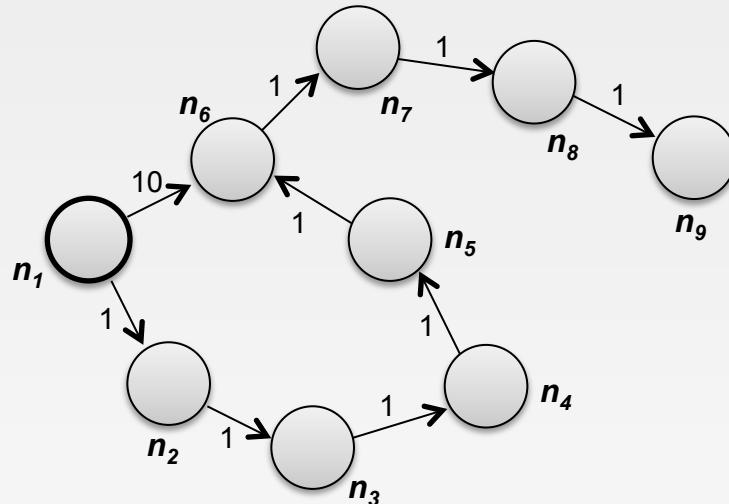
# Additional Complexities



- Assume that  $p$  is the current processed node
  - In the current iteration, we just “discovered” node  $r$  for the very first time.
  - We’ve already discovered the shortest distance to node  $p$ , and that the shortest distance to  $r$  so far goes through  $p$
  - Is  $s \rightarrow p \rightarrow r$  the shortest path from  $s$  to  $r$ ?
- The shortest path from source  $s$  to node  $r$  may go outside the current search frontier
  - It is possible that  $p \rightarrow q \rightarrow r$  is shorter than  $p \rightarrow r$ !
  - We will not find the shortest distance to  $r$  until the search frontier expands to cover  $q$ .

# How Many Iterations Are Needed?

- In the worst case, we might need as many iterations as there are nodes in the graph minus one
  - A sample graph that elicits worst-case behaviour for parallel breadth-first search.
  - Eight iterations are required to discover shortest distances to all nodes from  $n_1$ .



# Example (only distances)

## ■ Input file:

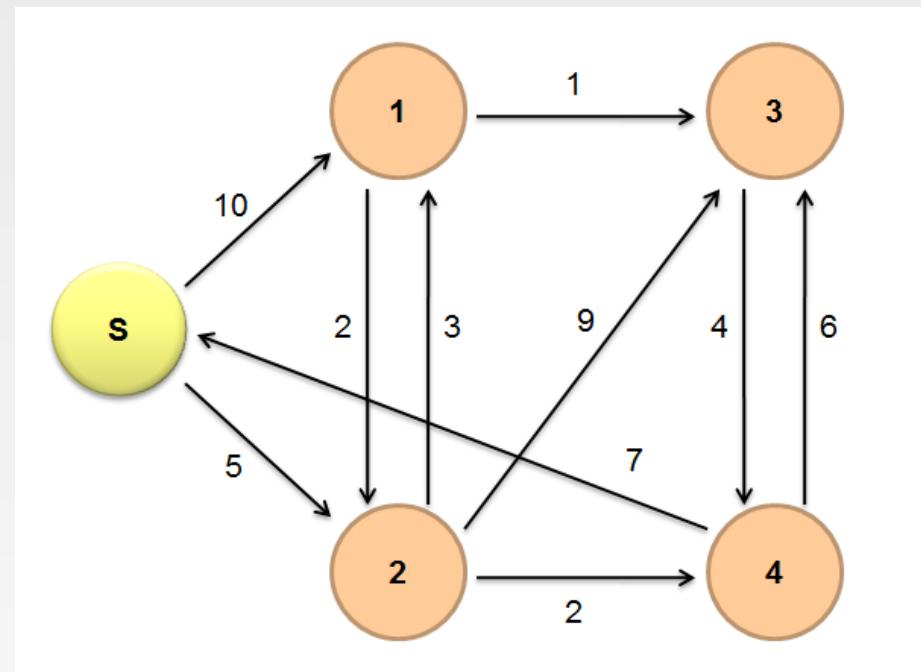
s --> 0 | n1: 10, n2: 5

n1 --> ∞ | n2: 2, n3:1

n2 --> ∞ | n1: 3, n3:9, n4:2

n3 --> ∞ | n4:4

n4 --> ∞ | s:7, n3:6



# Iteration 1

## ■ Map:

Read  $s \rightarrow 0 | n1: 10, n2: 5$

Emit:  $(n1, 10)$ ,  $(n2, 5)$ , and the adjacency list  $(s, n1: 10, n2: 5)$

*The other lists will also be read and emit, but they do not contribute, and thus ignored*

## ■ Reduce:

Receives:  $(n1, 10)$ ,  $(n2, 5)$ ,  $(s, <0, (n1: 10, n2: 5)>)$

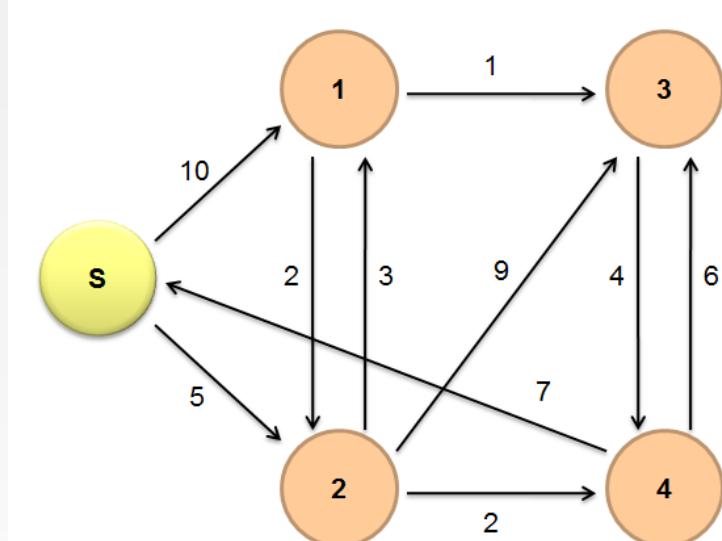
*The adjacency list of each node will also be received, ignored in example*

Emit:

$s \rightarrow 0 | n1: 10, n2: 5$

$n1 \rightarrow 10 | n2: 2, n3: 1$

$n2 \rightarrow 5 | n1: 3, n3: 9, n4: 2$



# Iteration 2

## ■ Map:

Read: n1 --> 10 | n2: 2, n3:1

Emit: (n2, 12), (n3, 11), (n1, <10, (n2: 2, n3:1)>)

Read: n2 --> 5 | n1: 3, n3:9, n4:2

Emit: (n1, 8), (n3, 14), (n4, 7), (n2, <5, (n1: 3, n3:9, n4:2)>)

*Ignore the processing of the other lists*

## ■ Reduce:

Receives: (n1, (8, <10, (n2: 2, n3:1)>)), (n2, (12, <5, n1: 3, n3:9, n4:2>)),  
(n3, (11, 14)), (n4, 7)

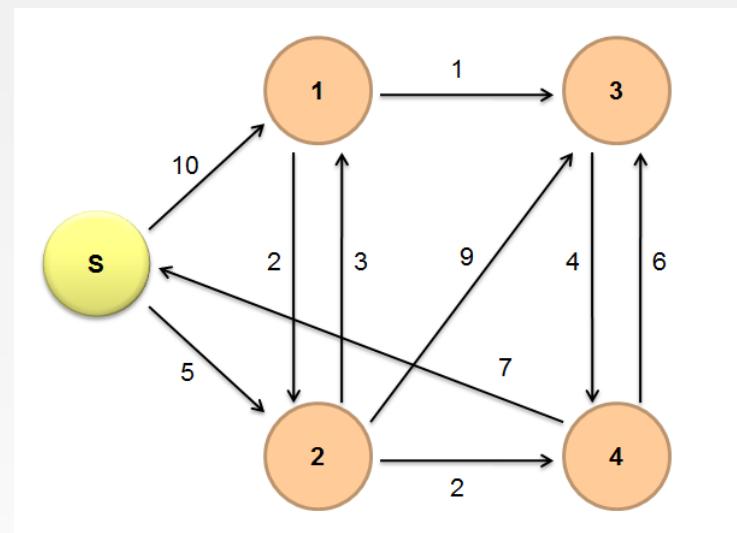
Emit:

n1 --> 8 | n2: 2, n3:1

n2 --> 5 | n1: 3, n3:9, n4:2

n3 --> 11 | n4:4

n4 --> 7 | s:7, n3:6



# Iteration 3

## ■ Map:

Read: n1 --> 8 | n2: 2, n3:1

Emit: (n2, 10), (n3, 9), (n1, <8, (n2: 2, n3:1)>)

Read: n2 --> 5 | n1: 3, n3:9, n4:2 (**Again!**)

Emit: (n1, 8), (n3, 14), (n4, 7), (n2, <5, (n1: 3, n3:9, n4:2)>)

Read: n3 --> 11 | n4:4

Emit: (n4, 15), (n3, <11, (n4:4)>)

Read: n4 --> 7 | s:7, n3:6

Emit: (s, 14), (n3, 13), (n4, <7, (s:7, n3:6)>)

## ■ Reduce:

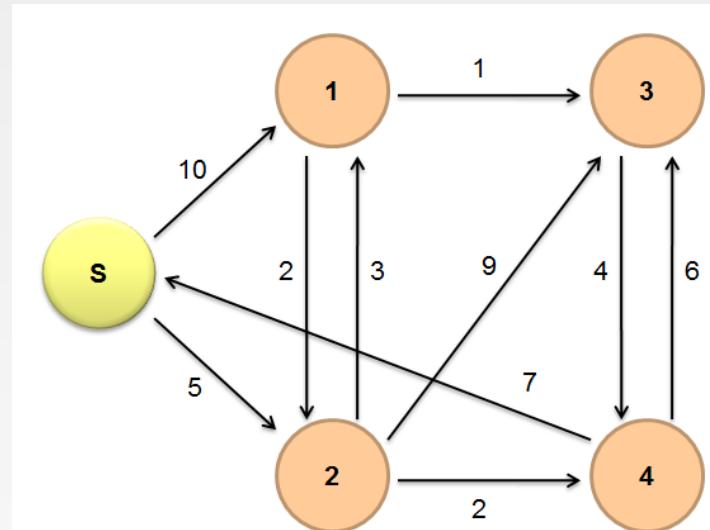
Emit:

n1 --> 8 | n2: 2, n3:1

n2 --> 5 | n1: 3, n3:9, n4:2

n3 --> 9 | n4:4

n4 --> 7 | s:7, n3:6



# Iteration 4

## ■ Map:

Read:  $n_1 \rightarrow 8 \mid n_2: 2, n_3: 1$  (**Again!**)

Emit:  $(n_2, 10), (n_3, 9), (n_1, <8, (n_2: 2, n_3: 1)>)$

Read:  $n_2 \rightarrow 5 \mid n_1: 3, n_3: 9, n_4: 2$  (**Again!**)

Emit:  $(n_1, 8), (n_3, 14), (n_4, 7), (n_2, <5, (n_1: 3, n_3: 9, n_4: 2)>)$

Read:  $n_3 \rightarrow 9 \mid n_4: 4$

Emit:  $(n_4, 13), (n_3, <9, (n_4: 4)>)$

Read:  $n_4 \rightarrow 7 \mid s: 7, n_3: 6$  (**Again!**)

Emit:  $(s, 14), (n_3, 13), (n_4, <7, (s: 7, n_3: 6)>)$

## ■ Reduce:

Emit:

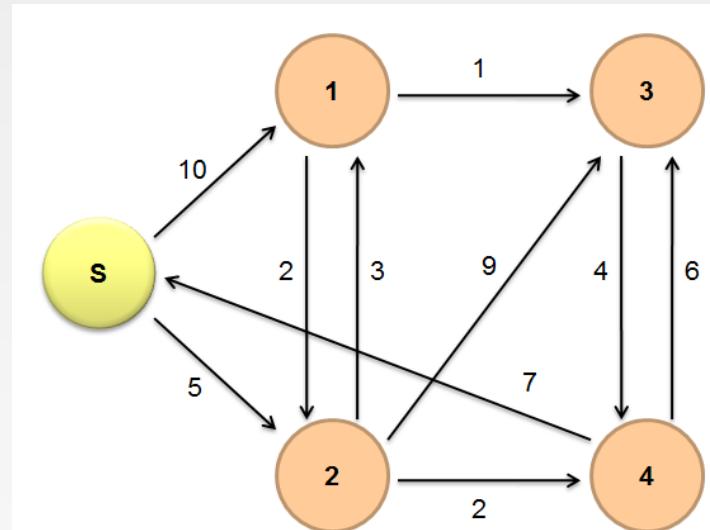
$n_1 \rightarrow 8 \mid n_2: 2, n_3: 1$

$n_2 \rightarrow 5 \mid n_1: 3, n_3: 9, n_4: 2$

$n_3 \rightarrow 9 \mid n_4: 4$

$n_4 \rightarrow 7 \mid s: 7, n_3: 6$

In order to avoid duplicated computations, you can use a status value to indicate whether the distance of the node has been modified in the previous iteration.



No updates. Terminate.

# Comparison to Dijkstra

- Dijkstra's algorithm is more efficient
  - At any step it only pursues edges from the minimum-cost path inside the frontier
- MapReduce explores all paths in parallel
  - Lots of “waste”
  - Useful work is only done at the “frontier”
- Why can't we do better using MapReduce?

# Graphs and MapReduce

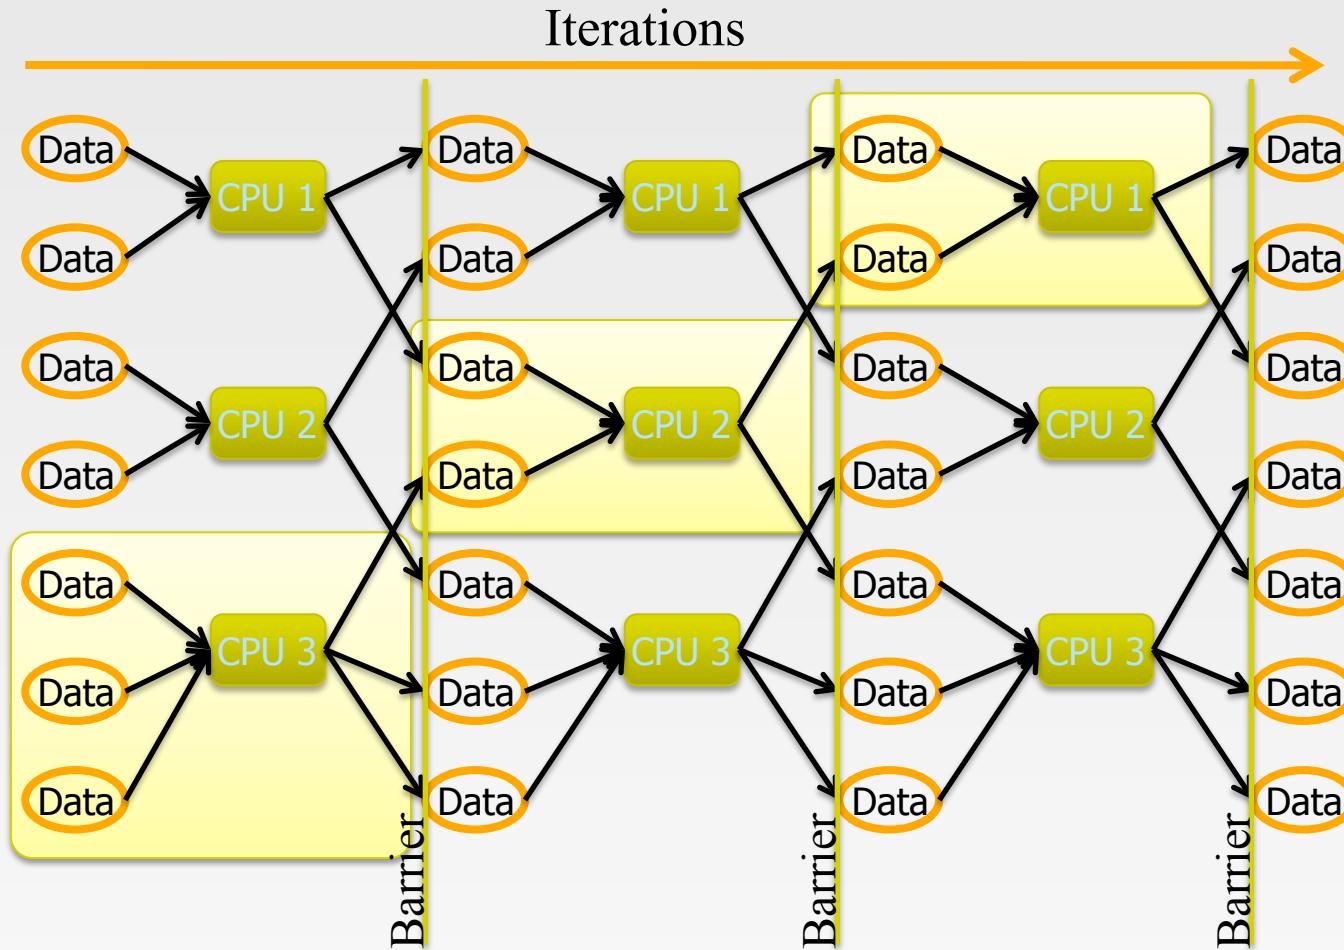
- Graph algorithms typically involve:
  - Performing computations at each node: based on node features, edge features, and local link structure
  - Propagating computations: “traversing” the graph
- Generic recipe:
  - Represent graphs as adjacency lists
  - Perform local computations in mapper
  - Pass along partial results via outlinks, keyed by destination node
  - Perform aggregation in reducer on inlinks to a node
  - Iterate until convergence: controlled by external “driver”
  - Don’t forget to pass the graph structure between iterations

# Issues with MapReduce on Graph Processing

- MapReduce Does not support iterative graph computations:
  - External driver. Huge I/O incurs
  - No mechanism to support global data structures that can be accessed and updated by all mappers and reducers
    - ▶ Passing information is only possible within the local graph structure – through adjacency list
    - ▶ Dijkstra's algorithm on a single machine: a global priority queue that guides the expansion of nodes
    - ▶ Dijkstra's algorithm in Hadoop, no such queue available. Do some “wasted” computation instead
- MapReduce algorithms are often impractical on large, dense graphs.
  - The amount of intermediate data generated is on the order of the number of edges.
  - For dense graphs, MapReduce running time would be dominated by copying intermediate data across the network.

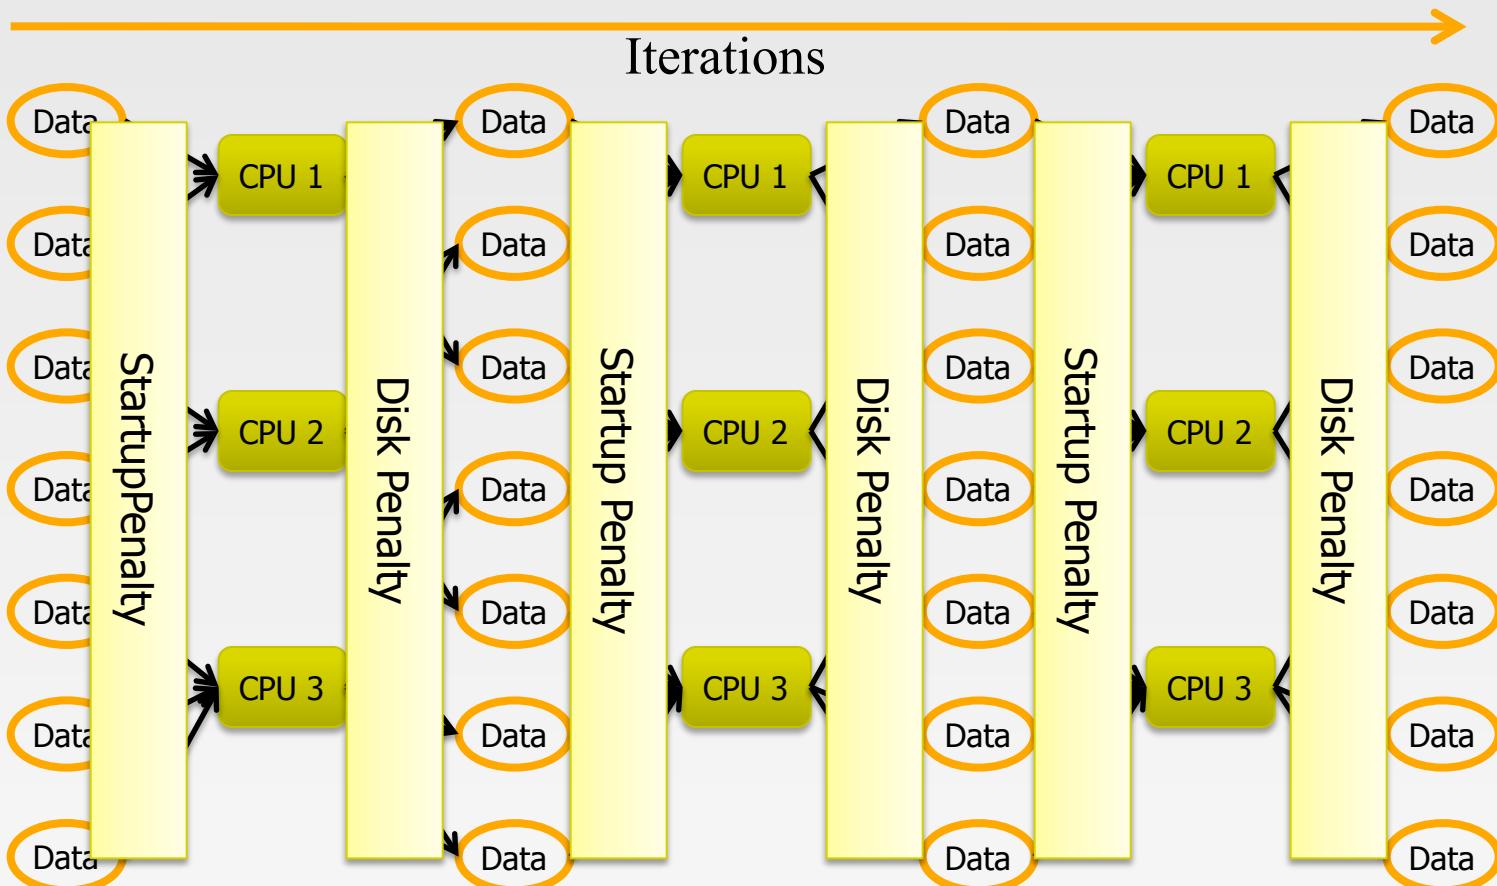
# Iterative MapReduce

- Only a subset of data needs computation:



# Iterative MapReduce

- System is not optimized for iteration:



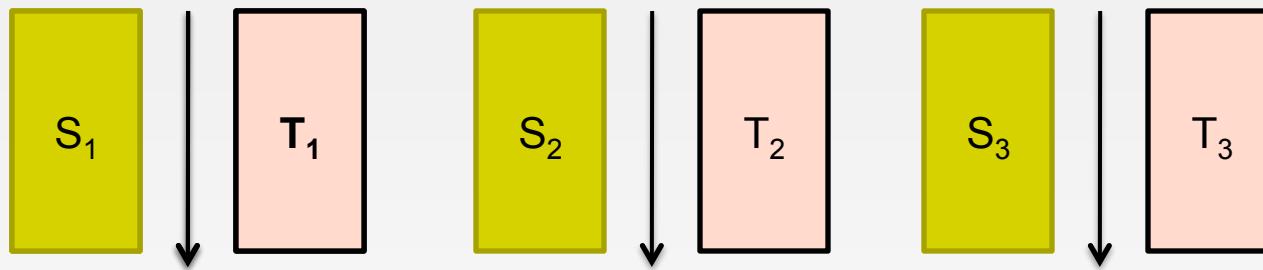
# Better Partitioning

- Default: hash partitioning
  - Randomly assign nodes to partitions
- Observation: many graphs exhibit local structure
  - E.g., communities in social networks
  - Better partitioning creates more opportunities for local aggregation
- Unfortunately, partitioning is **hard!**
  - Sometimes, chick-and-egg...
  - But cheap heuristics sometimes available
  - For webgraphs: range partition on domain-sorted URLs

# Schimmy Design Pattern

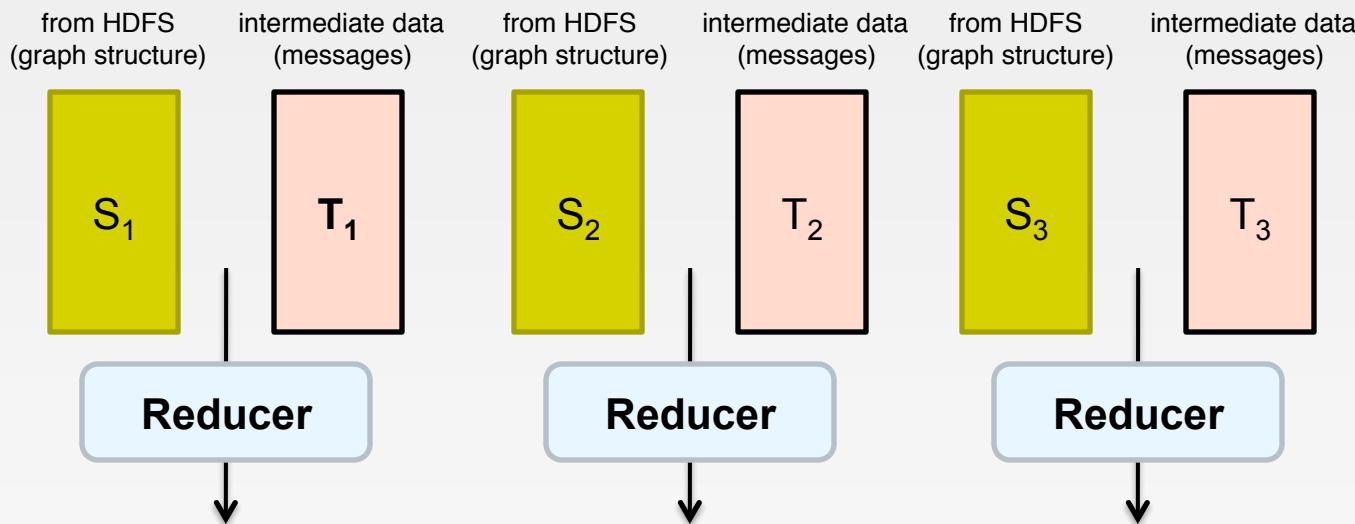
- Basic implementation contains two dataflows:
  - Messages (actual computations)
  - Graph structure (“bookkeeping”)
- Schimmy: separate the two dataflows, shuffle only the messages
  - Basic idea: merge join between graph structure and messages

both relationships/partitions/join keys consistently partitioned and sorted by join key



# Do the Schimmy!

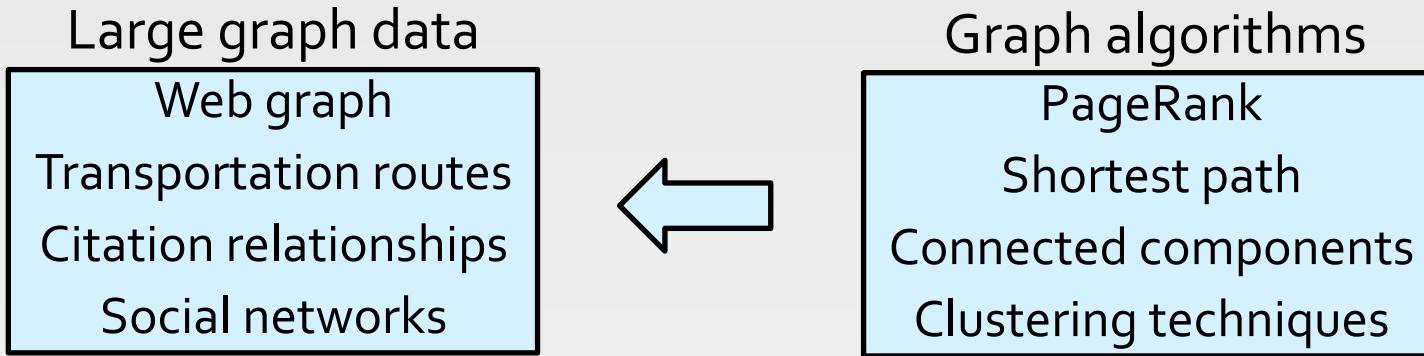
- Schimmy = reduce side parallel merge join between graph structure and messages
  - Consistent partitioning between input and intermediate data
  - Mappers emit only messages (actual computation)
  - Reducers read graph structure directly from HDFS



# **Introduction to Pregel**

# Motivation of Pregel

- Many practical computing problems concern large graphs



- Single computer graph library does not scale
- MapReduce is ill-suited for graph processing
  - Many iterations are needed for parallel graph processing
  - Materializations of intermediate results at every MapReduce iteration harm performance

# Pregel

- **Pregel**: A System for Large-Scale **Graph** Processing (Google) - Malewicz et al. SIGMOD 2010.
- Scalable and Fault-tolerant platform
- API with flexibility to express arbitrary algorithm
- Inspired by Valiant's Bulk Synchronous Parallel model
  - Leslie G. Valiant: A Bridging Model for Parallel Computation. Commun. ACM 33 (8): 103-111 (1990)
- Vertex centric computation (Think like a vertex)

# Bulk Synchronous Parallel Model (BSP)

analogous to MapReduce rounds

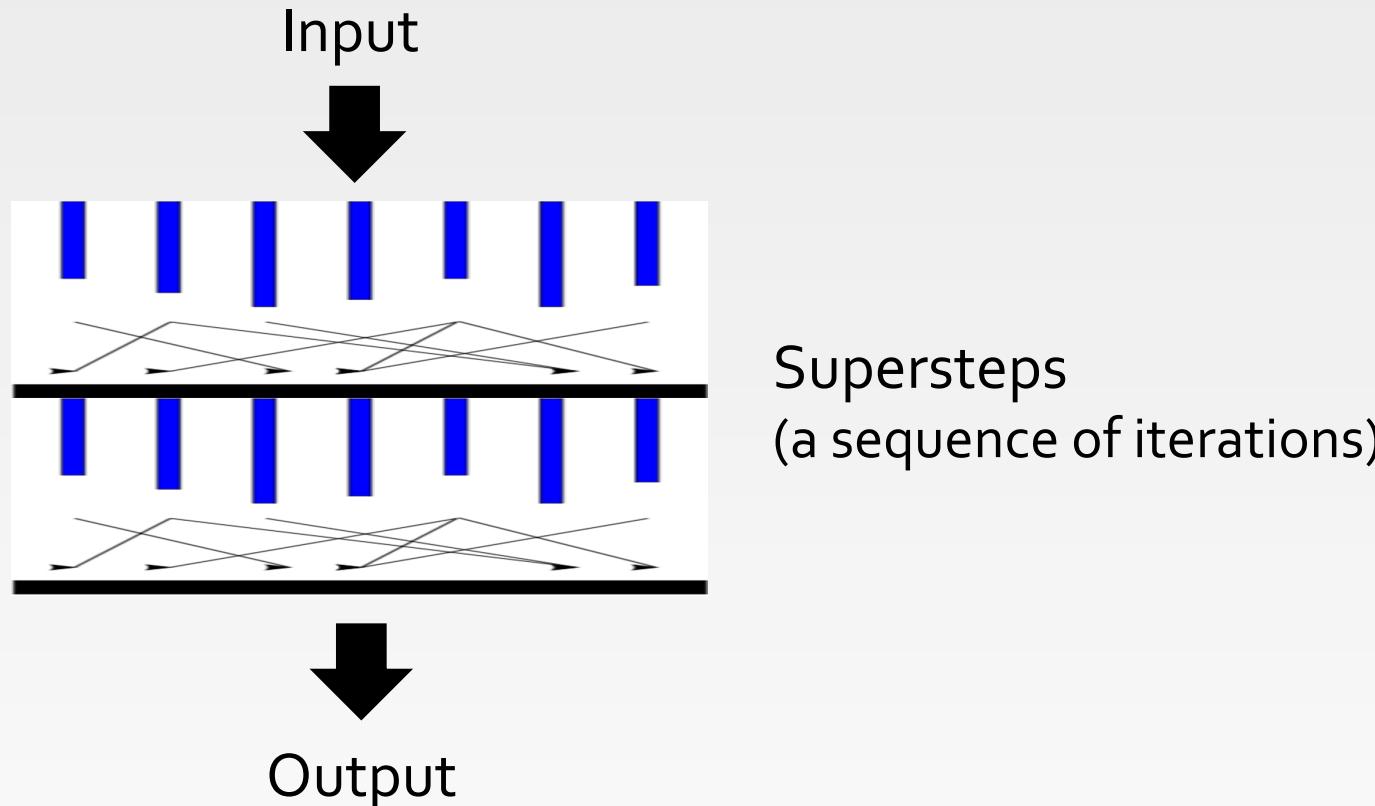
- Processing: a series of **supersteps**
- **Vertex**: computation is defined to run on each vertex
- Superstep S: *all vertices compute in parallel; each vertex v may*
  - receive **messages** sent to v from superstep S – 1;
  - perform some computation: modify its states and the states of its outgoing edges
  - Send **messages** to other vertices ( to be received in the next superstep)

Message passing

*Vertex-centric, message passing*

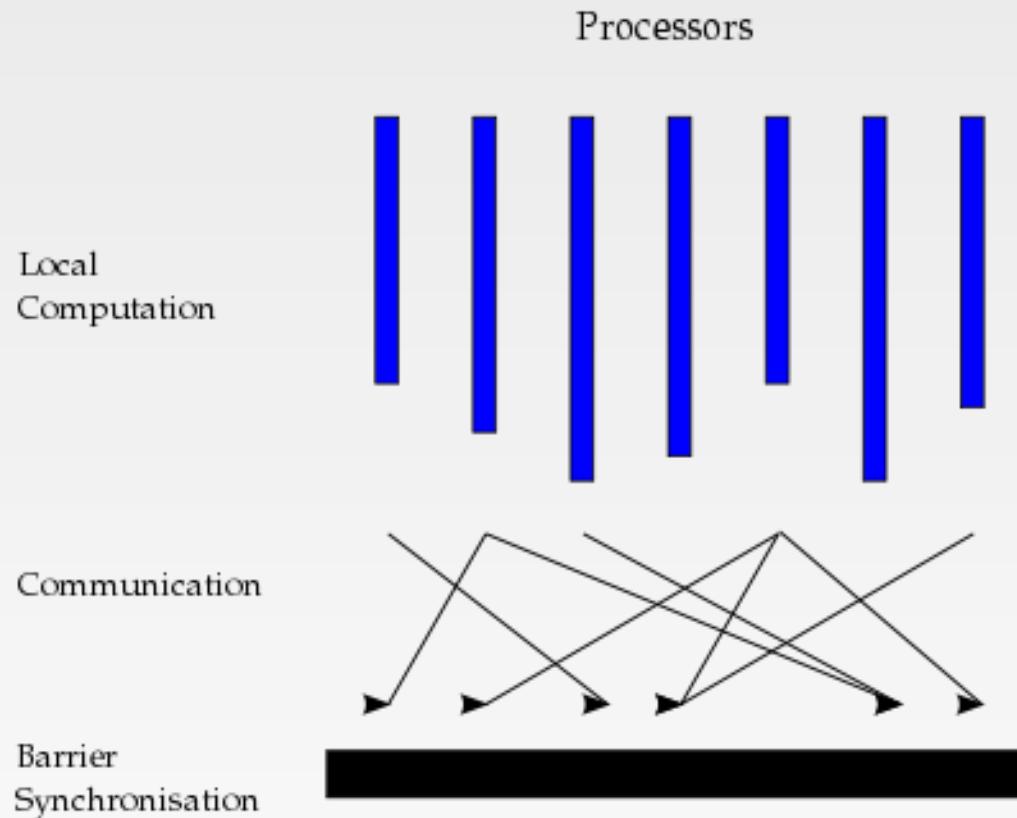
# Pregel Computation Model

- Based on Bulk Synchronous Parallel (BSP)
  - Computational units encoded in a directed graph
  - Computation proceeds in a series of supersteps
  - Message passing architecture



# Pregel Computation Model (Cont')

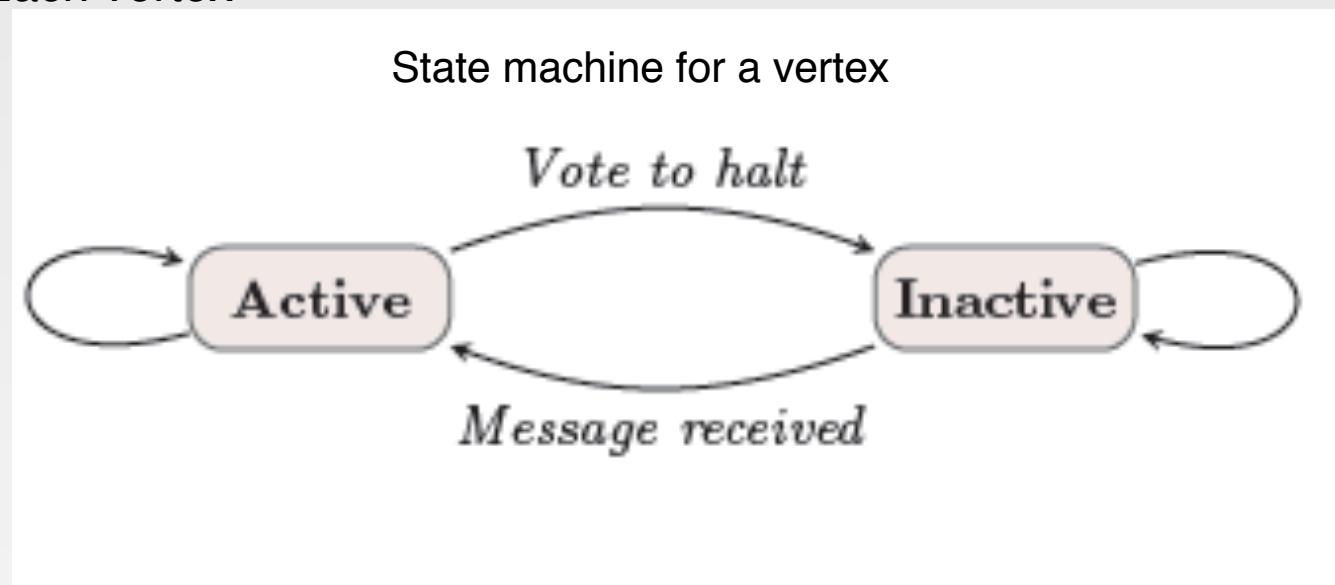
- Concurrent computation and Communication need not be ordered in time
- Communication through message passing



Source: [http://en.wikipedia.org/wiki/Bulk\\_synchronous\\_parallel](http://en.wikipedia.org/wiki/Bulk_synchronous_parallel)

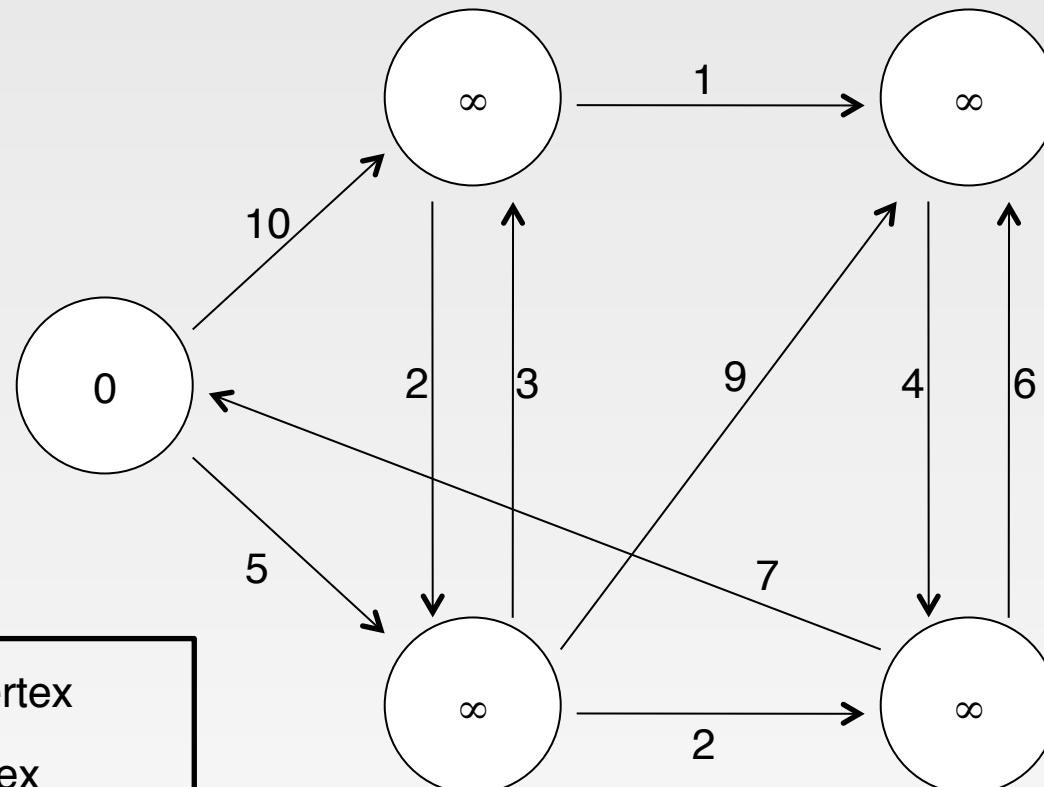
# Pregel Computation Model (Cont')

- Superstep: the vertices compute in parallel
  - Each vertex



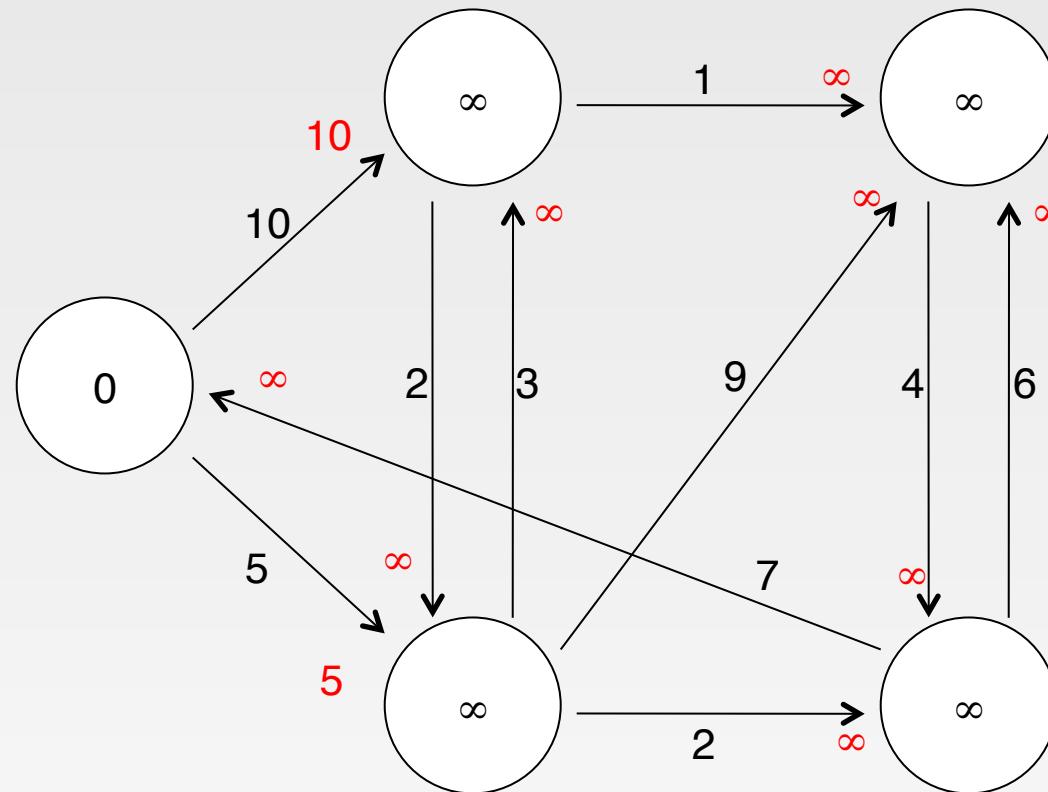
- Termination condition
  - ▶ All vertices are simultaneously inactive
  - ▶ A vertex can choose to deactivate itself
  - ▶ Is “woken up” if new messages received

# Example: SSSP – Parallel BFS in Pregel

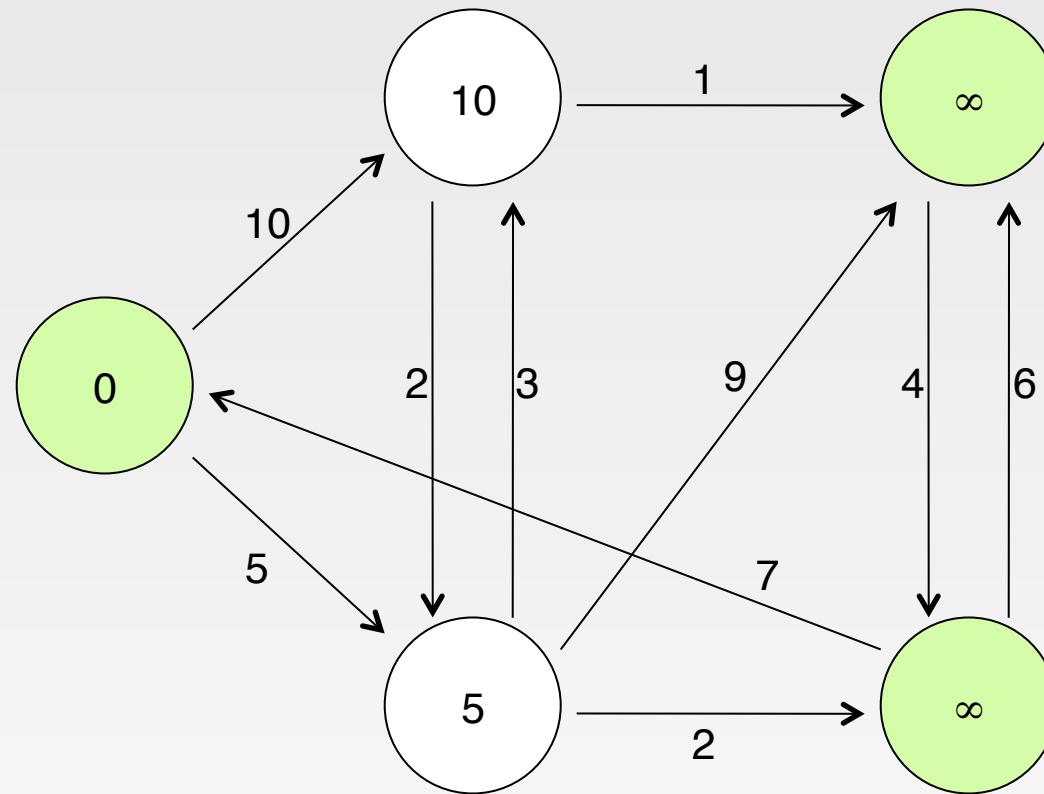


- Inactive Vertex
- Active Vertex
- Edge weight
- ✗ Message

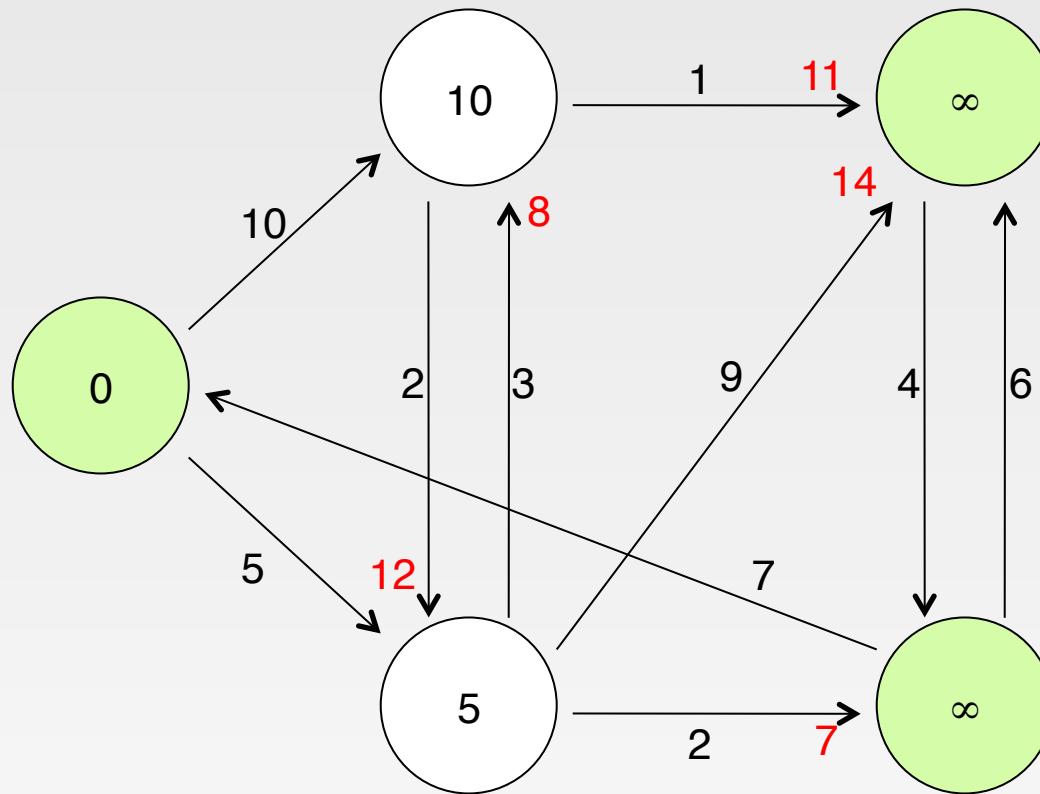
# Example: SSSP – Parallel BFS in Pregel



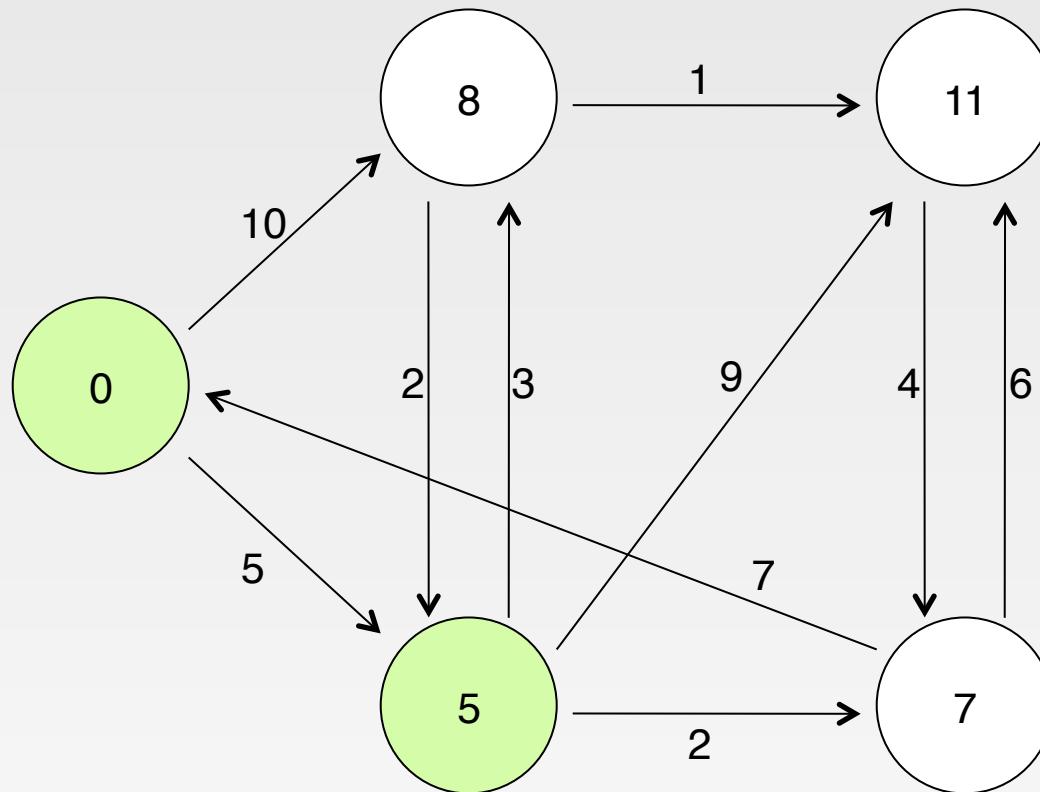
# Example: SSSP – Parallel BFS in Pregel



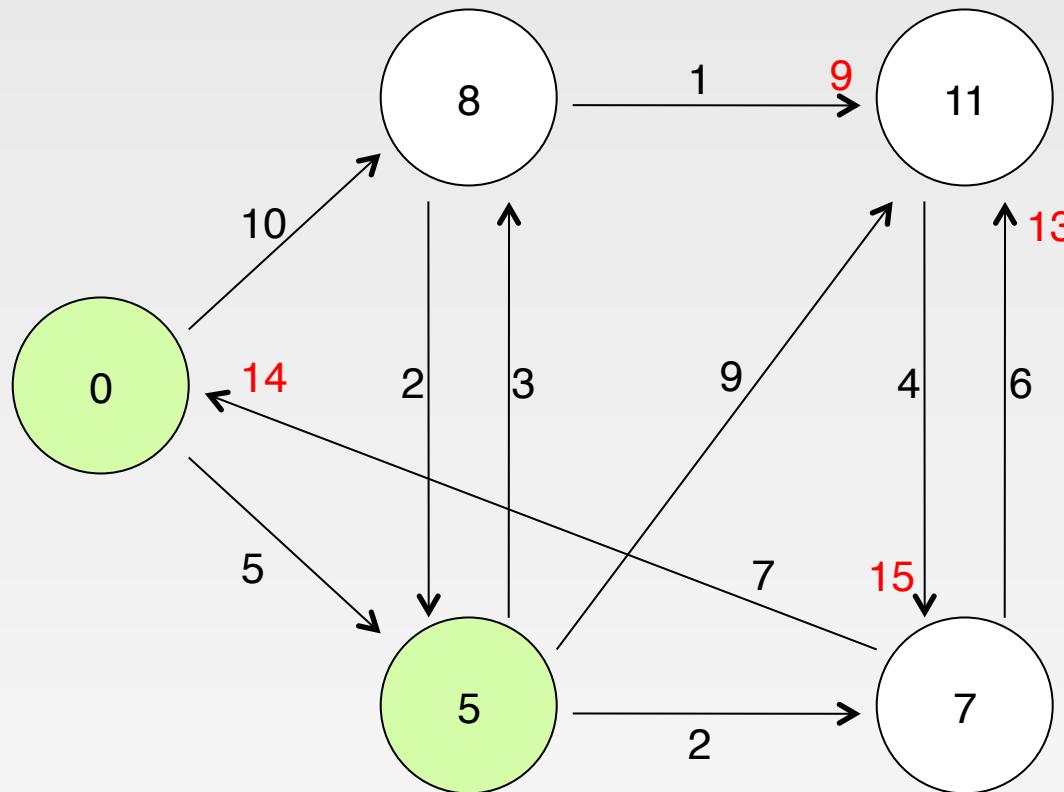
# Example: SSSP – Parallel BFS in Pregel



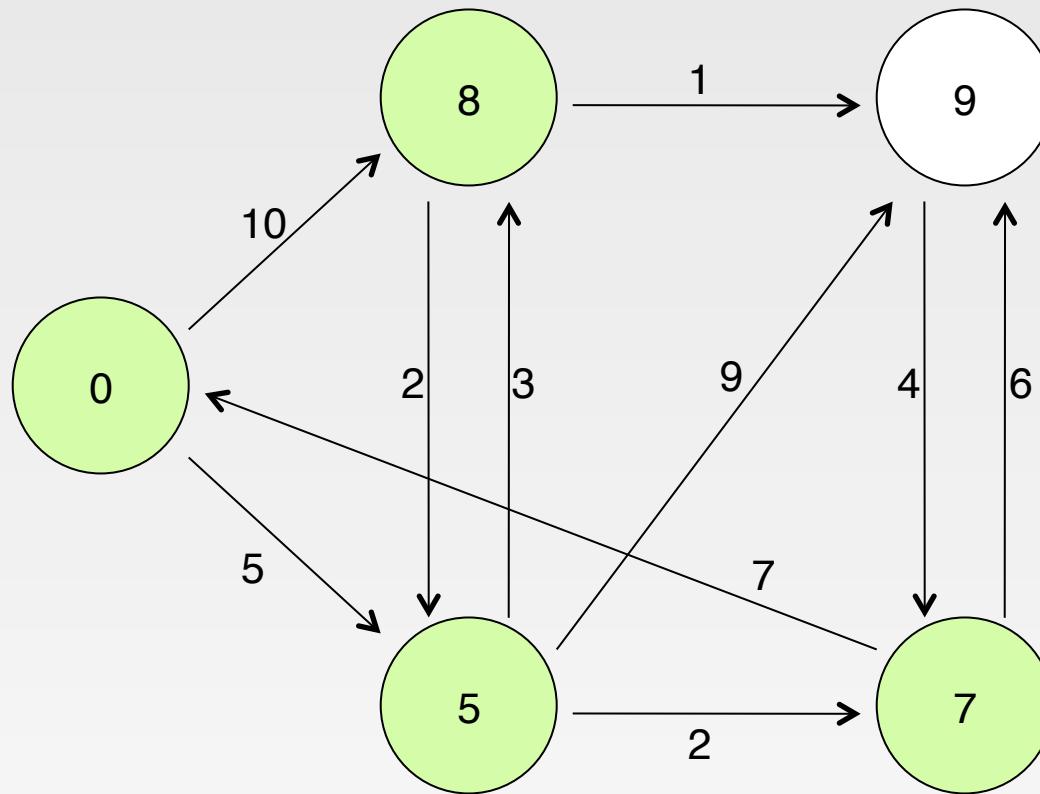
# Example: SSSP – Parallel BFS in Pregel



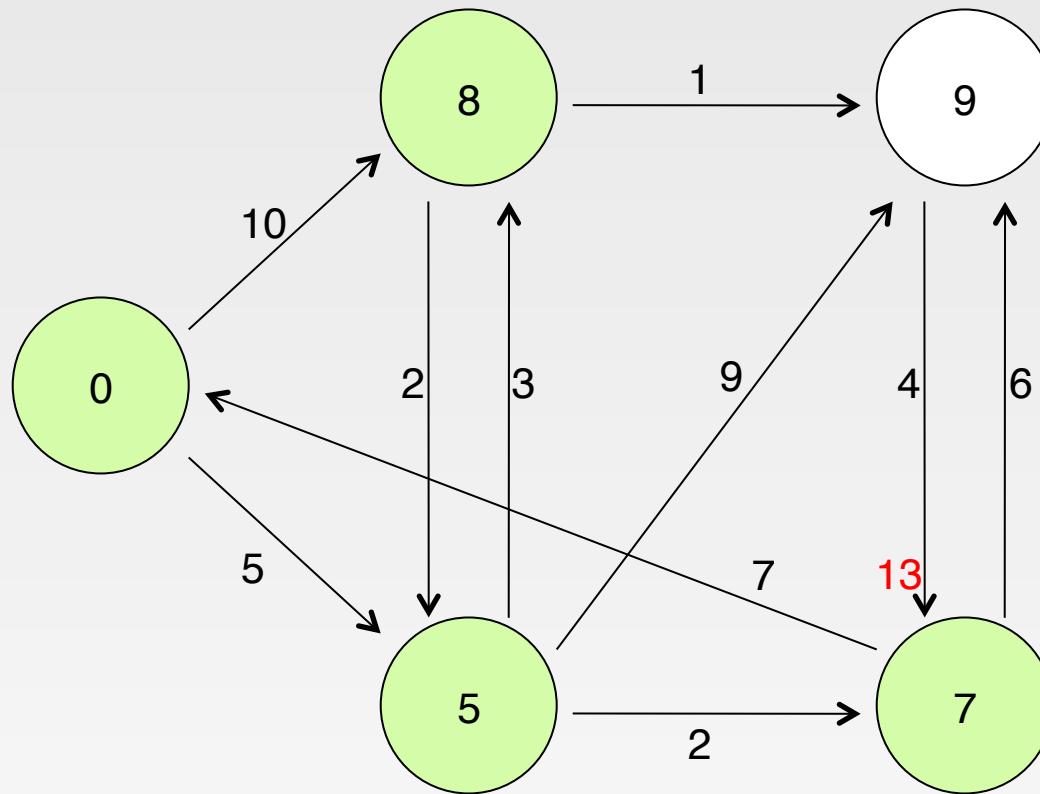
# Example: SSSP – Parallel BFS in Pregel



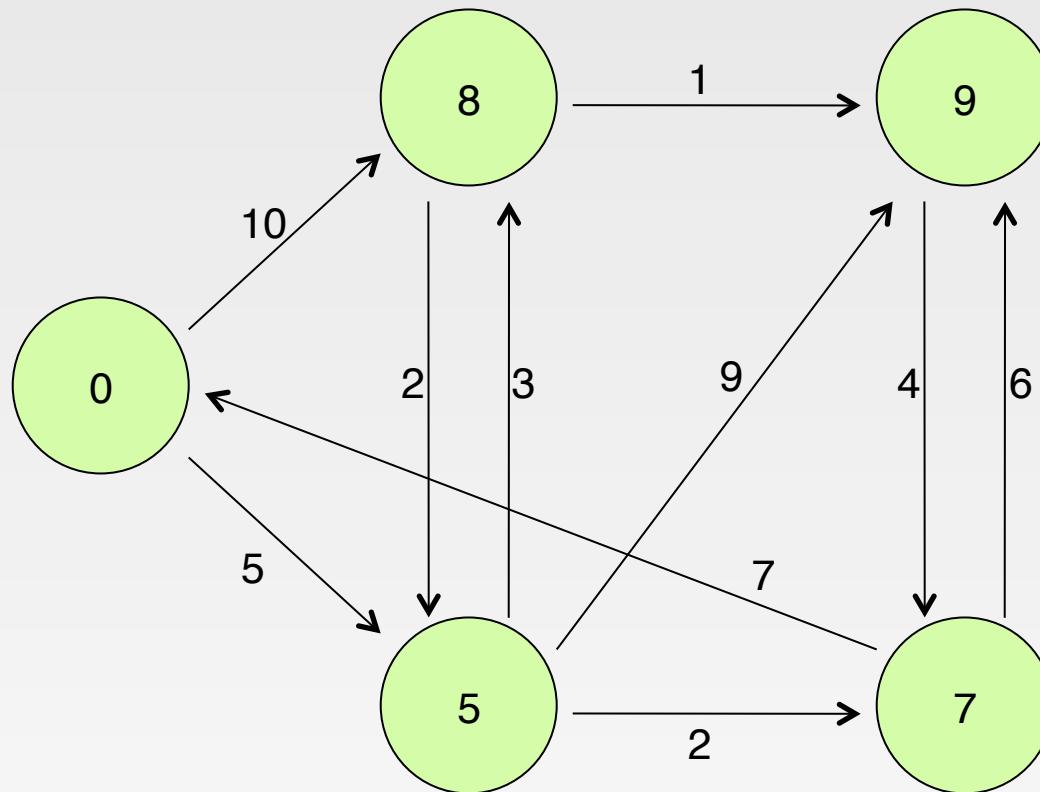
# Example: SSSP – Parallel BFS in Pregel



# Example: SSSP – Parallel BFS in Pregel



# Example: SSSP – Parallel BFS in Pregel



# Differences from MapReduce

- Graph algorithms can be written as a series of chained MapReduce jobs
- Pregel
  - Keeps vertices & edges on the machine that performs computation
  - Uses network transfers only for messages
- MapReduce
  - Passes the entire state of the graph from one stage to the next
  - Needs to coordinate the steps of a chained MapReduce

# Writing a Pregel Program (C++)

## ■ Subclassing the predefined **Vertex** class

```
template <typename VertexValue,  
          typename EdgeValue,  
          typename MessageValue>  
class Vertex {  
public:  
    virtual void Compute(MessageIterator* msgs) = 0;  
  
    const string& vertex_id() const;  
    int64 superstep() const;  
  
    const VertexValue& GetValue();  
    VertexValue* MutableValue();  
    OutEdgeIterator GetOutEdgeIterator();  
    void SendMessageTo(const string& dest_vertex,  
                      const MessageValue& message);  
    void VoteToHalt();  
};
```

Override this!

in msgs

Modify vertex value

out msg

# Pregel: SSSP (C++)

```
class ShortestPathVertex : public Vertex<int, int, int>
{
    void Compute(MessageIterator* msgs) { aggregation
        int mindist = IsSource(vertex_id()) ? 0 : INF;

        for (; !msgs->Done(); msgs->Next())
            mindist = min(mindist, msgs->Value()); // Messages: distances to u
        if (mindist < GetValue())
            *MutableValue() = mindist; // MutableValue: the current distance
        OutEdgeIterator iter = GetOutEdgeIterator();
        for (; !iter.Done(); iter.Next())
            SendMessageTo(iter.Target(), mindist + iter.GetValue()); // Pass revised distance to its neighbors
    }
    VoteToHalt();
}
```

Refer to the current node as  $u$

aggregation

Messages: distances to  $u$

MutableValue: the current distance

Pass revised distance to its neighbors

# More Tools on Big Graph Processing

## ■ Graph databases: Storage and Basic Operators

- [http://en.wikipedia.org/wiki/Graph\\_database](http://en.wikipedia.org/wiki/Graph_database)
- Neo4j (an open source graph database)
- InfiniteGraph
- VertexDB
- ... ...

## ■ Distributed Graph Processing (mostly in-memory-only)

- Google's Pregel (vertex centered computation)
- Giraph (Apache)
- GraphX (Spark)
- GraphLab
- ... ...

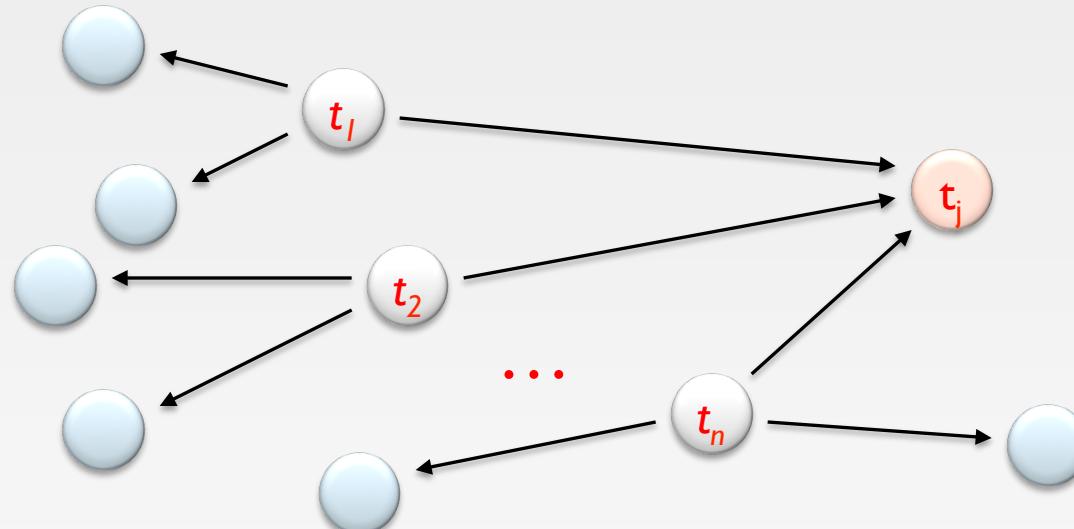
# References

- Chapter 5. Data-Intensive Text Processing with MapReduce
- Chapter 5. Mining of Massive Datasets.

# **End of Chapter 5**

# PageRank Review

- Given page  $t_j$  with in-coming neighbors  $t_1, \dots, t_n$ , where
    - $d_i$  is the out-degree of  $t_i$
    - $\beta$  is the teleport probability
    - $N$  is the total number of nodes in the graph
    - $r_{j\downarrow} = \sum_{i \rightarrow j} r_{i\downarrow} / d_i + (1 - \beta) 1/N$



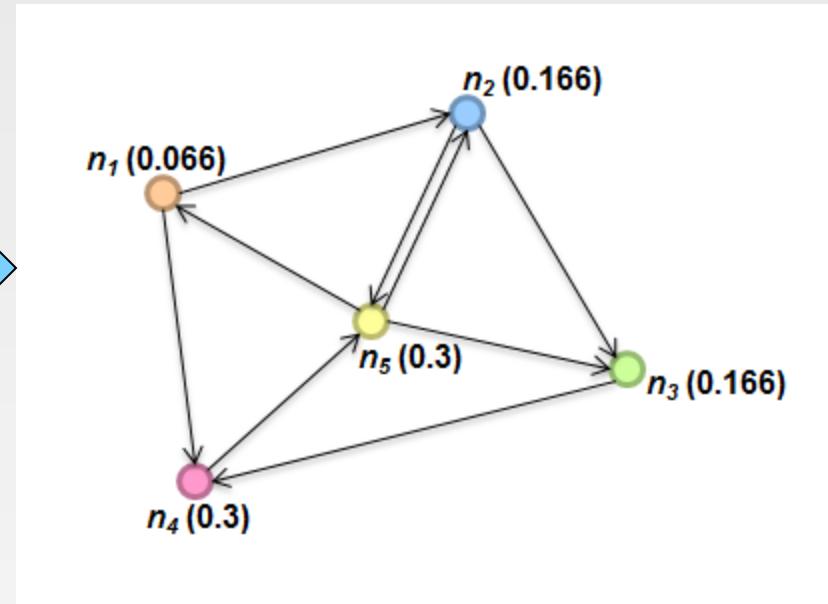
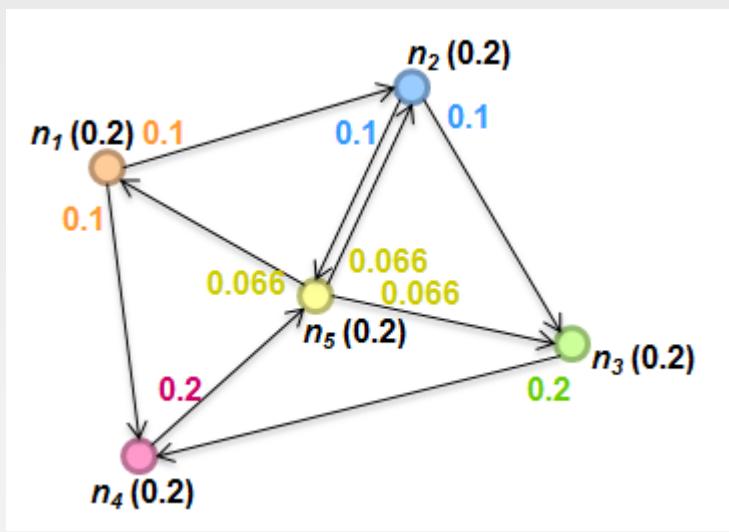
# Computing PageRank

- Properties of PageRank
  - Can be computed iteratively
  - Effects at each iteration are local
- Sketch of algorithm:
  - Start with seed  $r_i$  values
  - Each page distributes  $r_i$  “credit” to all pages it links to
  - Each target page  $t_j$  adds up “credit” from multiple in-bound links to compute  $r_j$
  - Iterate until values converge

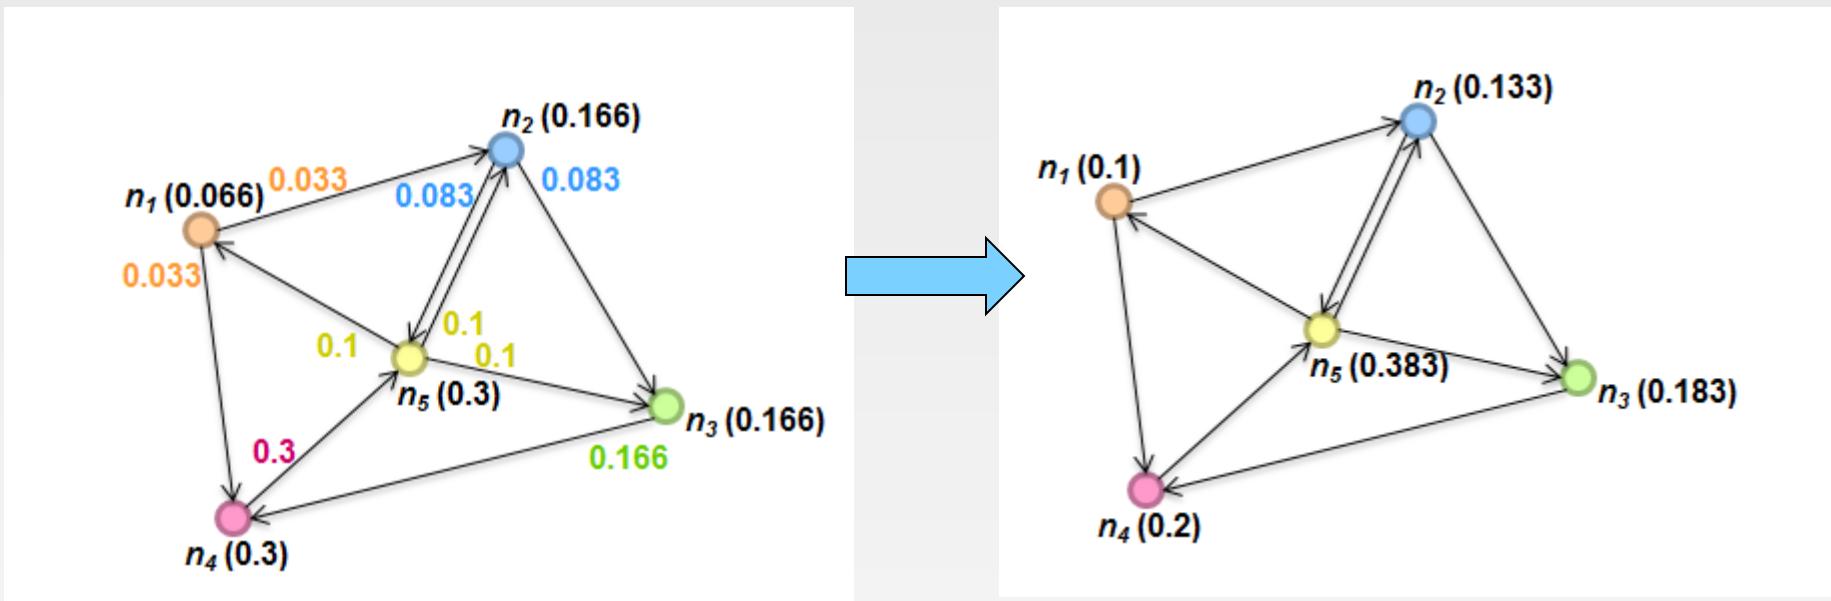
# Simplified PageRank

- First, tackle the simple case:
  - No teleport
  - No dangling nodes (dead ends)
- Then, factor in these complexities...
  - How to deal with the teleport probability?
  - How to deal with dangling nodes?

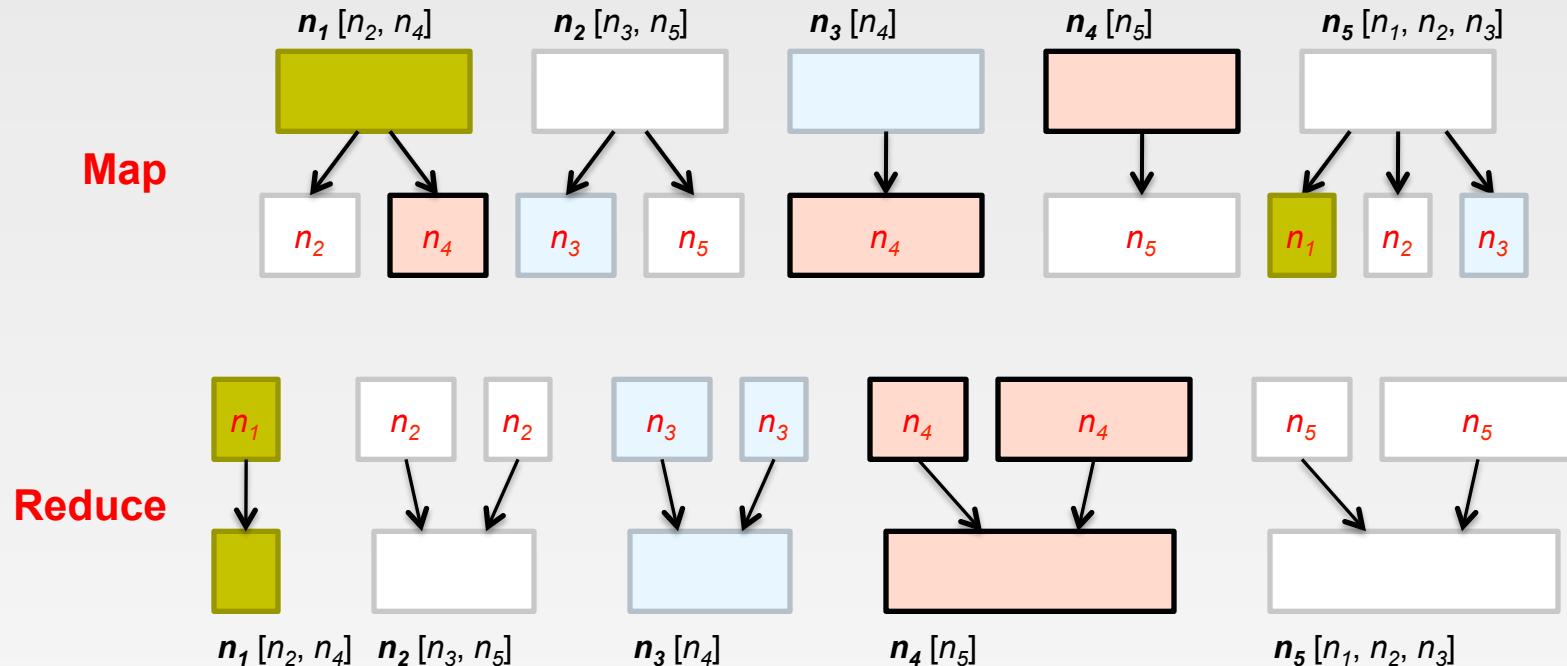
# Sample PageRank Iteration (1)



# Sample PageRank Iteration (2)



# PageRank in MapReduce (One Iteration)



# PageRank Pseudo-Code

```
1: class MAPPER
2:   method MAP(nid  $n$ , node  $N$ )
3:      $p \leftarrow N.\text{PAGERANK}/|N.\text{ADJACENCYLIST}|$ 
4:     EMIT(nid  $n, N$ )                                ▷ Pass along graph structure
5:     for all nodeid  $m \in N.\text{ADJACENCYLIST}$  do
6:       EMIT(nid  $m, p$ )                            ▷ Pass PageRank mass to neighbors
7:
8: class REDUCER
9:   method REDUCE(nid  $m, [p_1, p_2, \dots]$ )
10:     $M \leftarrow \emptyset$ 
11:    for all  $p \in \text{counts} [p_1, p_2, \dots]$  do
12:      if IsNODE( $p$ ) then
13:         $M \leftarrow p$                                 ▷ Recover graph structure
14:      else
15:         $s \leftarrow s + p$                           ▷ Sums incoming PageRank contributions
16:     $M.\text{PAGERANK} \leftarrow s$ 
17:    EMIT(nid  $m, \text{node } M$ )
```

# PageRank vs. BFS

PageRank

BFS

Map

$r_i/d_i$

$d+w$

Reduce

sum

min

# PageRank in Pregel

- Superstep 0: Value of each vertex is  $1/\text{NumVertices}()$

```
virtual void Compute(MessageIterator* msgs) {
    if (superstep() >= 1) {
        double sum = 0;
        for (; !msgs->done(); msgs->Next())
            sum += msgs->Value();
        *MutableValue() = 0.15 + 0.85 * sum;
    }
    if (supersteps() < 30) {
        const int64 n = GetOutEdgeIterator().size();
        SendMessageToAllNeighbors(GetValue() / n);
    } else {
        VoteToHalt();
    }
}
```