

# **COMP9313: Big Data Management**

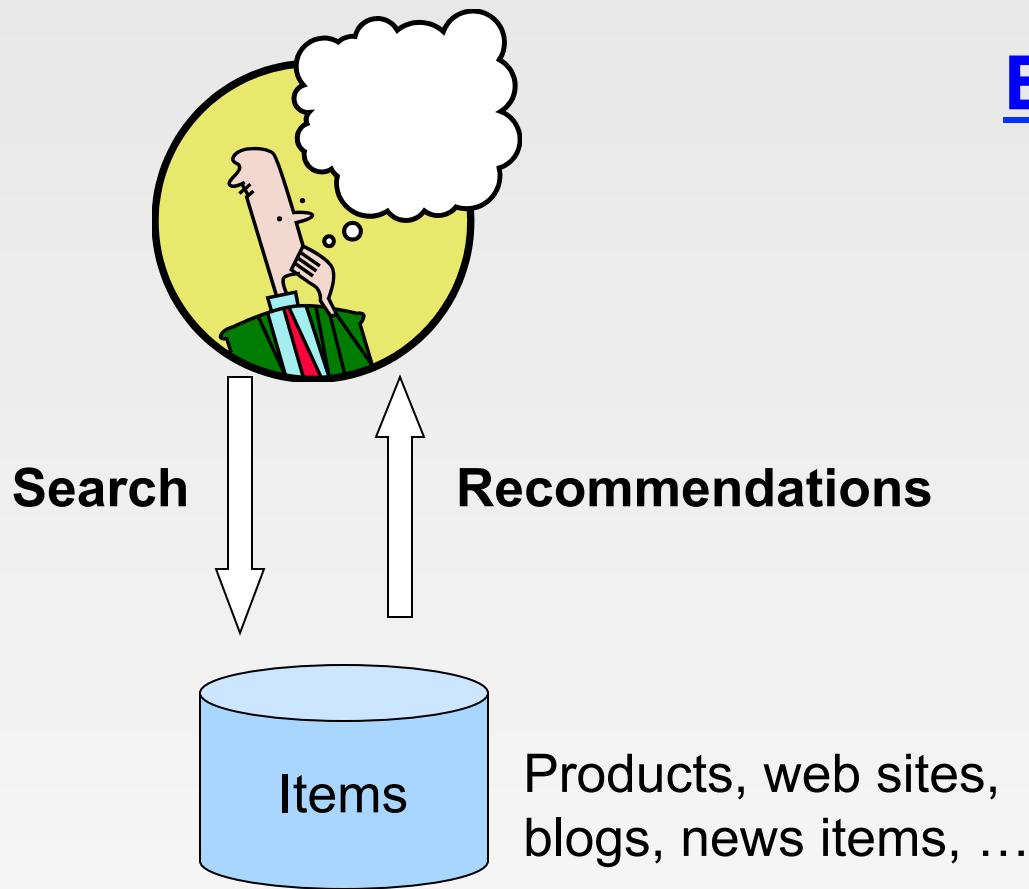


**Lecturer: Xin Cao**

**Course web site:** <http://www.cse.unsw.edu.au/~cs9313/>

# **Chapter 11: Recommender Systems**

# Recommendations



## Examples:

amazon.com.  
PANDORA®

StumbleUpon

del.icio.us

P  
PANDORA®

NETFLIX

movieLens  
helping you find the *right* movies

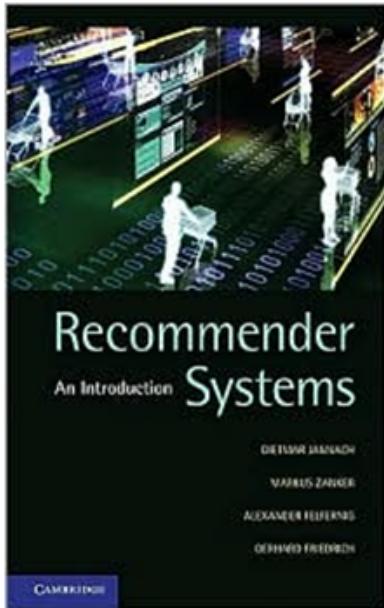
last.fm™  
the social music revolution

Google™  
News

You Tube

XBOX  
LIVE

# Recommender Systems

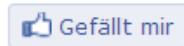


## Recommender Systems: An Introduction

by [Dietmar Jannach](#), [Markus Zanker](#), [Alexander Felfernig](#), [Gerhard Friedrich](#)

### AVERAGE CUSTOMER RATING:

( Be the first to review )



Registrieren, um sehen zu können, was  
deinen Freunden gefällt.

### FORMAT:

Hardcover

NOOKbook (eBook) - not available

[Tell the publisher you want this in NOOKbook format](#)

### NEW FROM BN.COM

\$65.00 List Price

**\$52.00** Online Price  
(You Save 20%)

**Add to Cart**

### NEW & USED FROM OUR

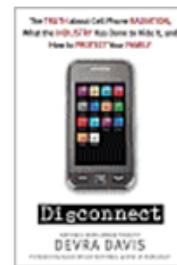
New starting at **\$56.46** (You S

Used starting at **\$51.98** (You S

**See All Prices**

### Table of Contents

### Customers who bought this also bought



# Recommender Systems

## Application areas

You may also like



Jack & Jones  
JAMIE - Polo shirt - orange  
£21.00  
Free delivery & returns

Recommended for You

RECOMMENDED APPS

- LinkedIn Popular with Foursquare users Free
- Yahoo! Messenger Plug-in Popular with Yahoo! Messenger users Free
- Firefox Beta Popular with Firefox users Free
- Windows Live Hotmail Popular with Hotmail users Mobi4All Tech Free
- Stars Live Wallpaper Popular with Blooming Night Live Wallpaper users Free

### Related hotels...



Hotel 41

★★★★★ 1,170 Reviews  
London, England

Show Prices

### Jobs you may be interested in Beta

Email Alerts | See More »



Technical Sales Manager - Europe

Thermal Transfer Products - Home office



Senior Program Manager (f/m)

Johnson Controls - Germany-NW-Burscheid

### You may also like



★★★★★ (109)



★★★★★ (53)



★★★★★ (33)

Picasa™ -Webalben

Startseite Meine Fotos Erkunden Hochladen

Empfohlene Fotos Alle anzeigen

# Why using Recommender Systems?

## ■ Value for the customer

- Find things that are interesting
- Narrow down the set of choices
- Help me explore the space of options
- Discover new things
- Entertainment
- ...

## ■ Value for the provider

- Additional and probably unique personalized service for the customer
- Increase trust and customer loyalty
- Increase sales, click trough rates, conversion etc.
- Opportunities for promotion, persuasion
- Obtain more knowledge about customers
- ...

# Real-world Check

## ■ Myths from industry

- Amazon.com generates X percent of their sales through the recommendation lists ( $30 < X < 70$ )
- Netflix (DVD rental and movie streaming) generates X percent of their sales through the recommendation lists ( $30 < X < 70$ )

# Recommender systems

- RS seen as a function
- Given:
  - User model (e.g. ratings, preferences, demographics, situational context)
  - Items (with or without description of item characteristics)
- Find:
  - Relevance score. Used for ranking.
- Finally:
  - Recommend items that are assumed to be relevant
- But:
  - Remember that relevance might be context-dependent
  - Characteristics of the list itself might be important (diversity)

# Formal Model

- $X$  = set of Customers
- $S$  = set of Items
- Utility function  $u: X \times S \rightarrow R$ 
  - $R$  = set of ratings
  - $R$  is a totally ordered set
  - e.g., 0-5 stars, real number in  $[0,1]$
- Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David			0.4	

# Key Problems

- Gathering “known” ratings for matrix
  - How to collect the data in the utility matrix
- Extrapolate unknown ratings from the known ones
  - Mainly interested in high unknown ratings
    - ▶ We are not interested in knowing what you don’t like but what you like
- Evaluating extrapolation methods
  - How to measure success/performance of recommendation methods

# Gathering Ratings

## ■ Explicit

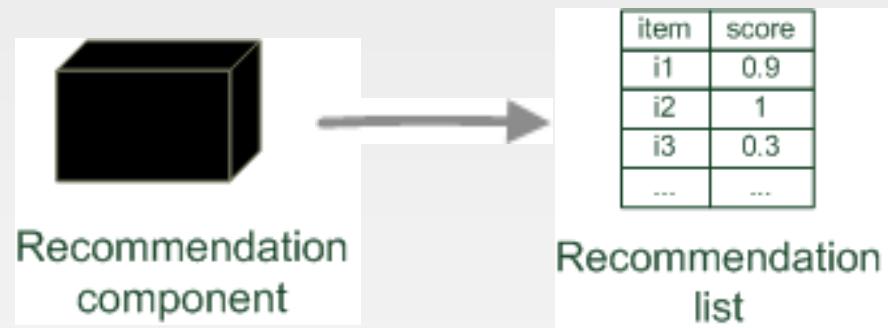
- Ask people to rate items
- Doesn't work well in practice – people can't be bothered

## ■ Implicit

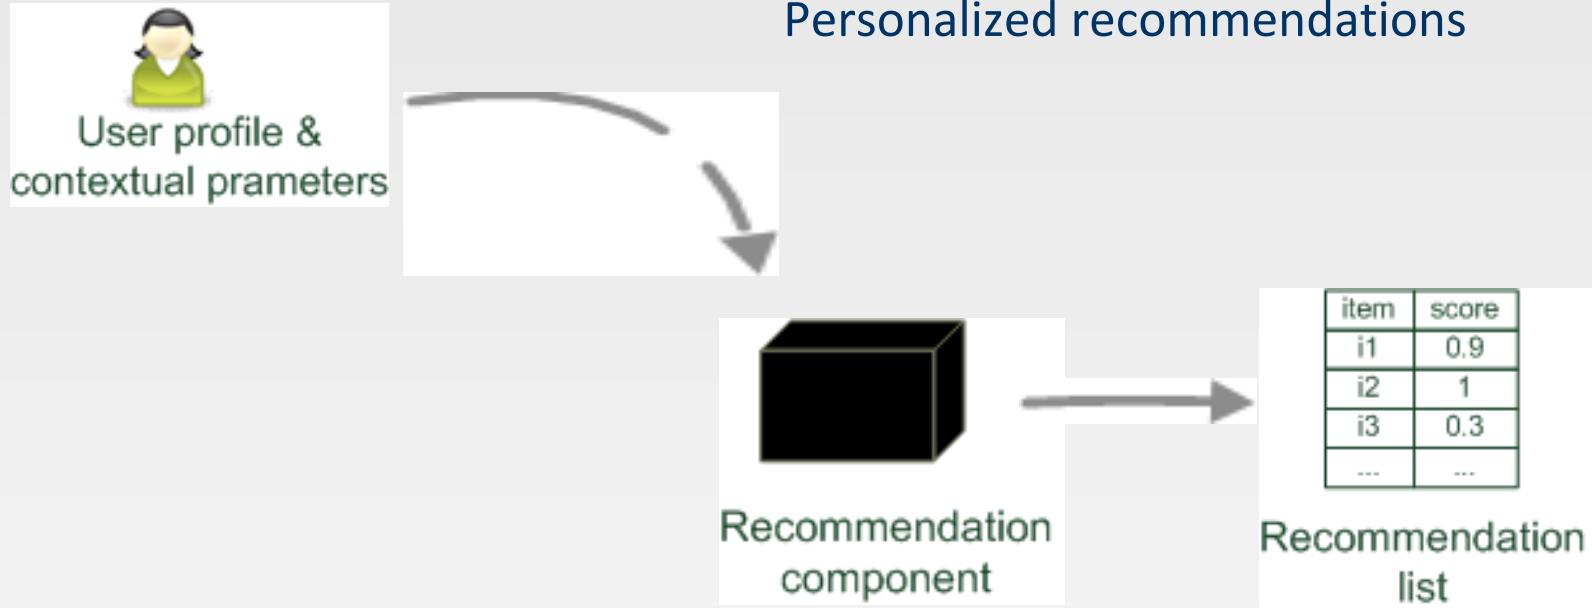
- Learn ratings from user actions
  - ▶ E.g., purchase implies high rating
- What about low ratings?

# Paradigms of recommender systems

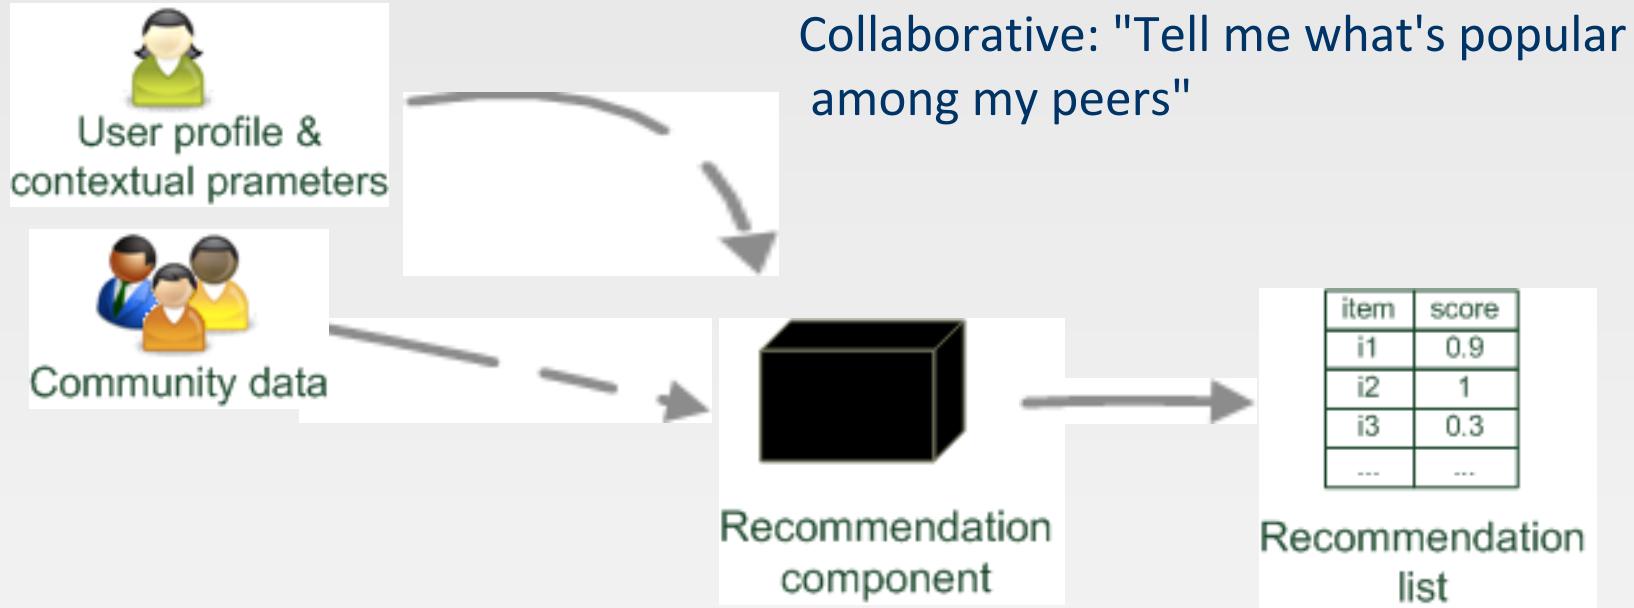
Recommender systems reduce information overload by estimating relevance



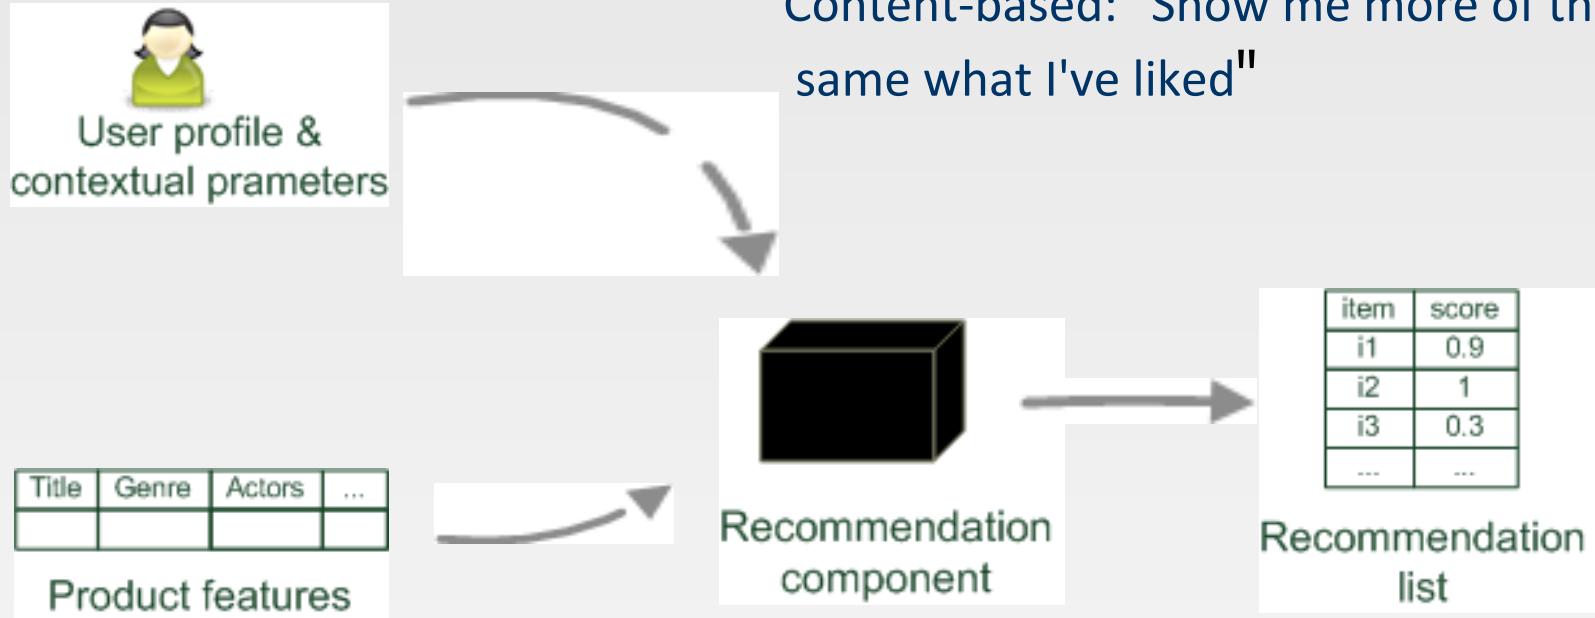
# Paradigms of recommender systems



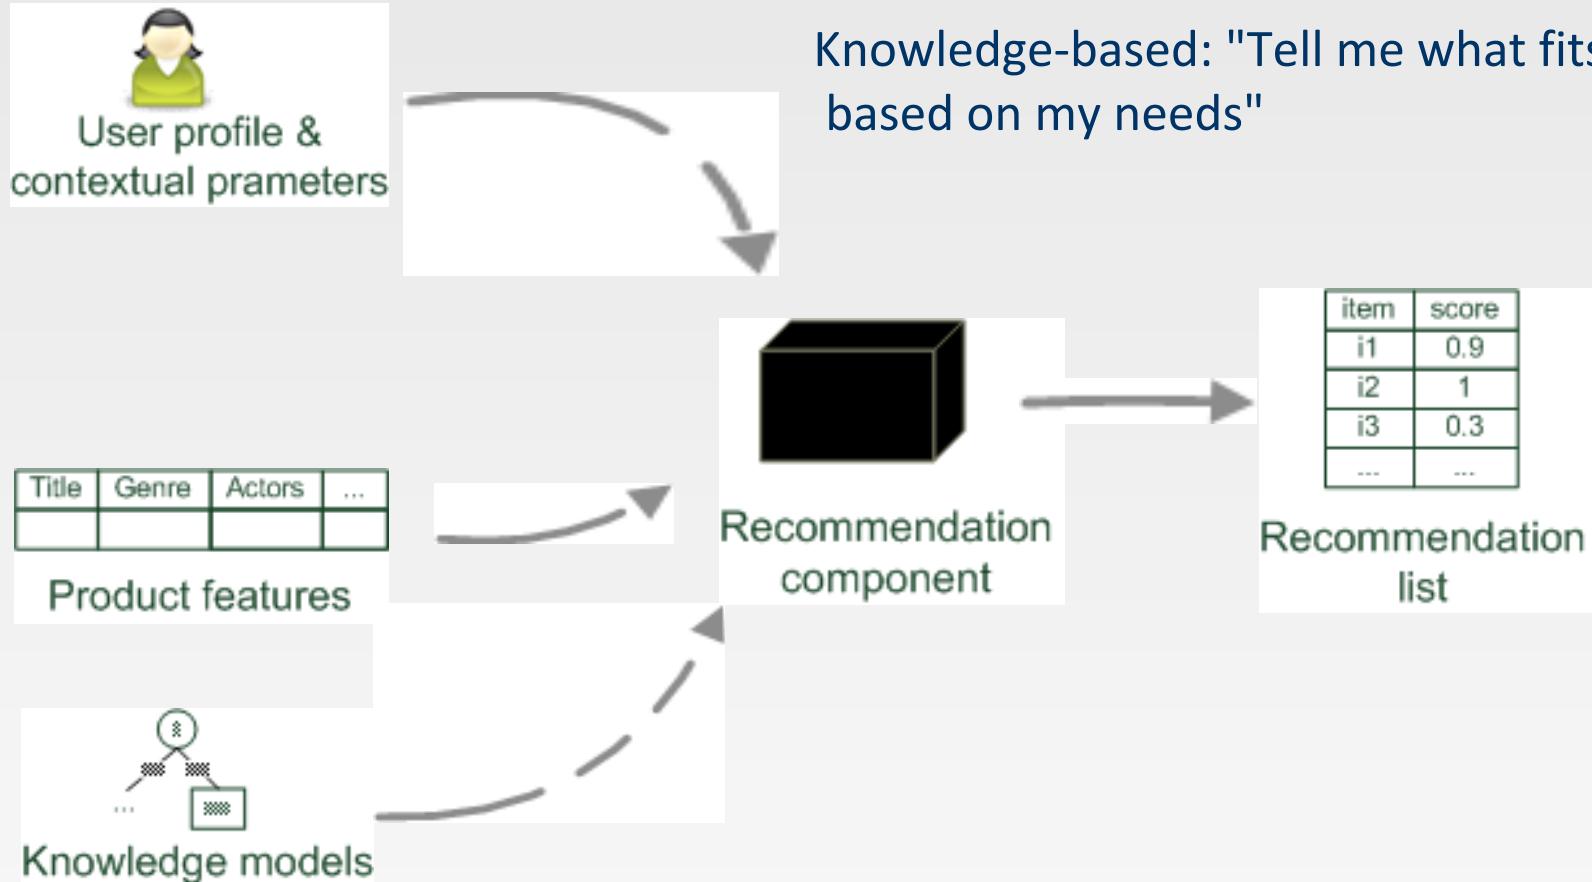
# Paradigms of recommender systems



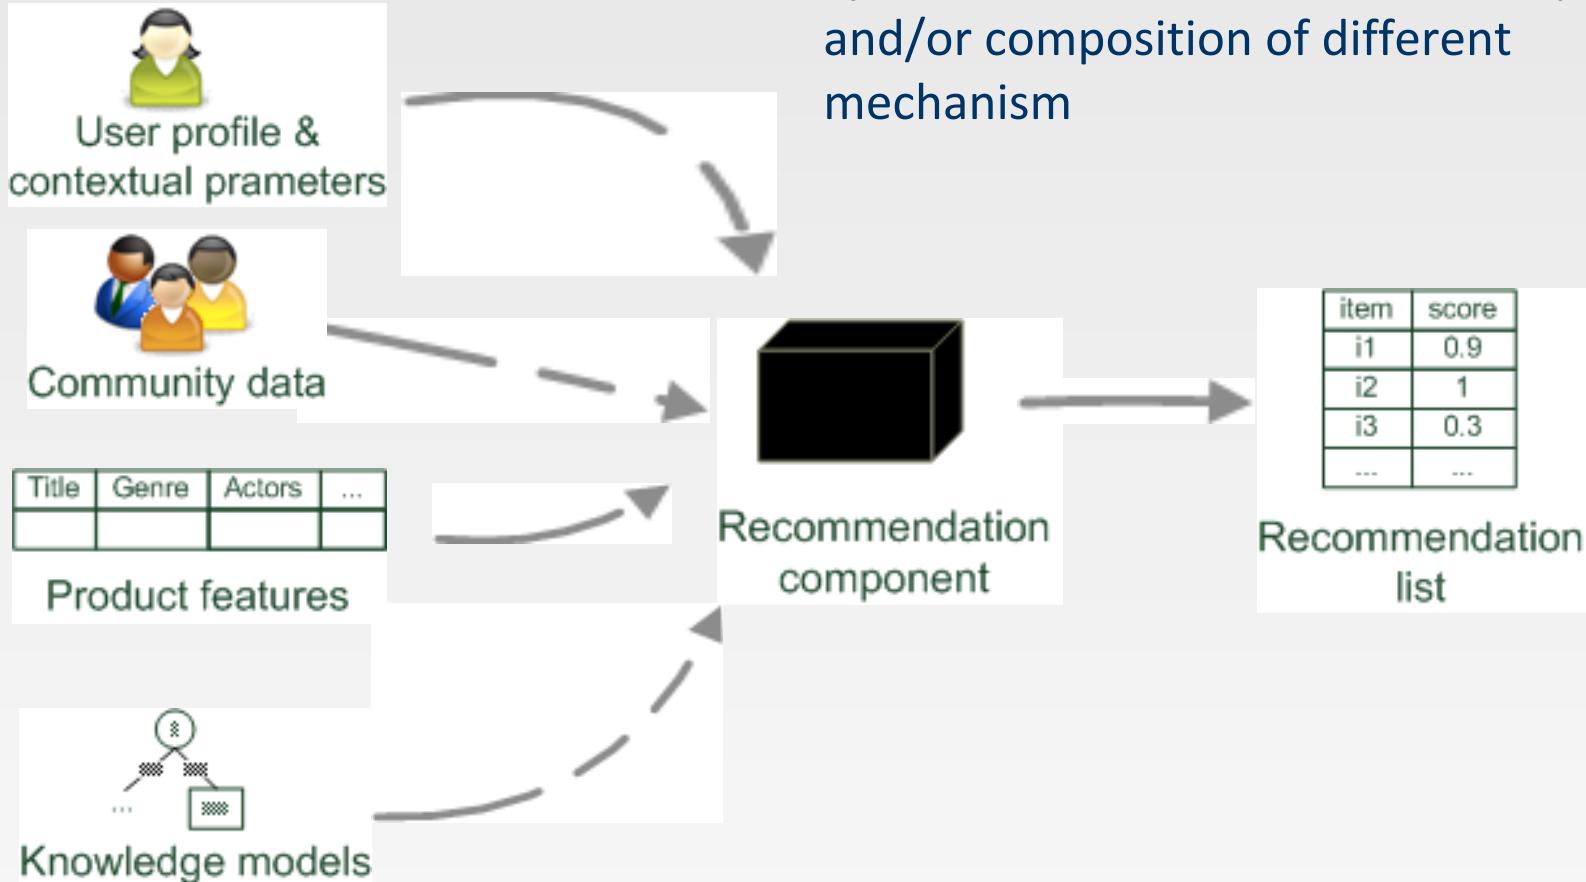
# Paradigms of recommender systems



# Paradigms of recommender systems



# Paradigms of recommender systems



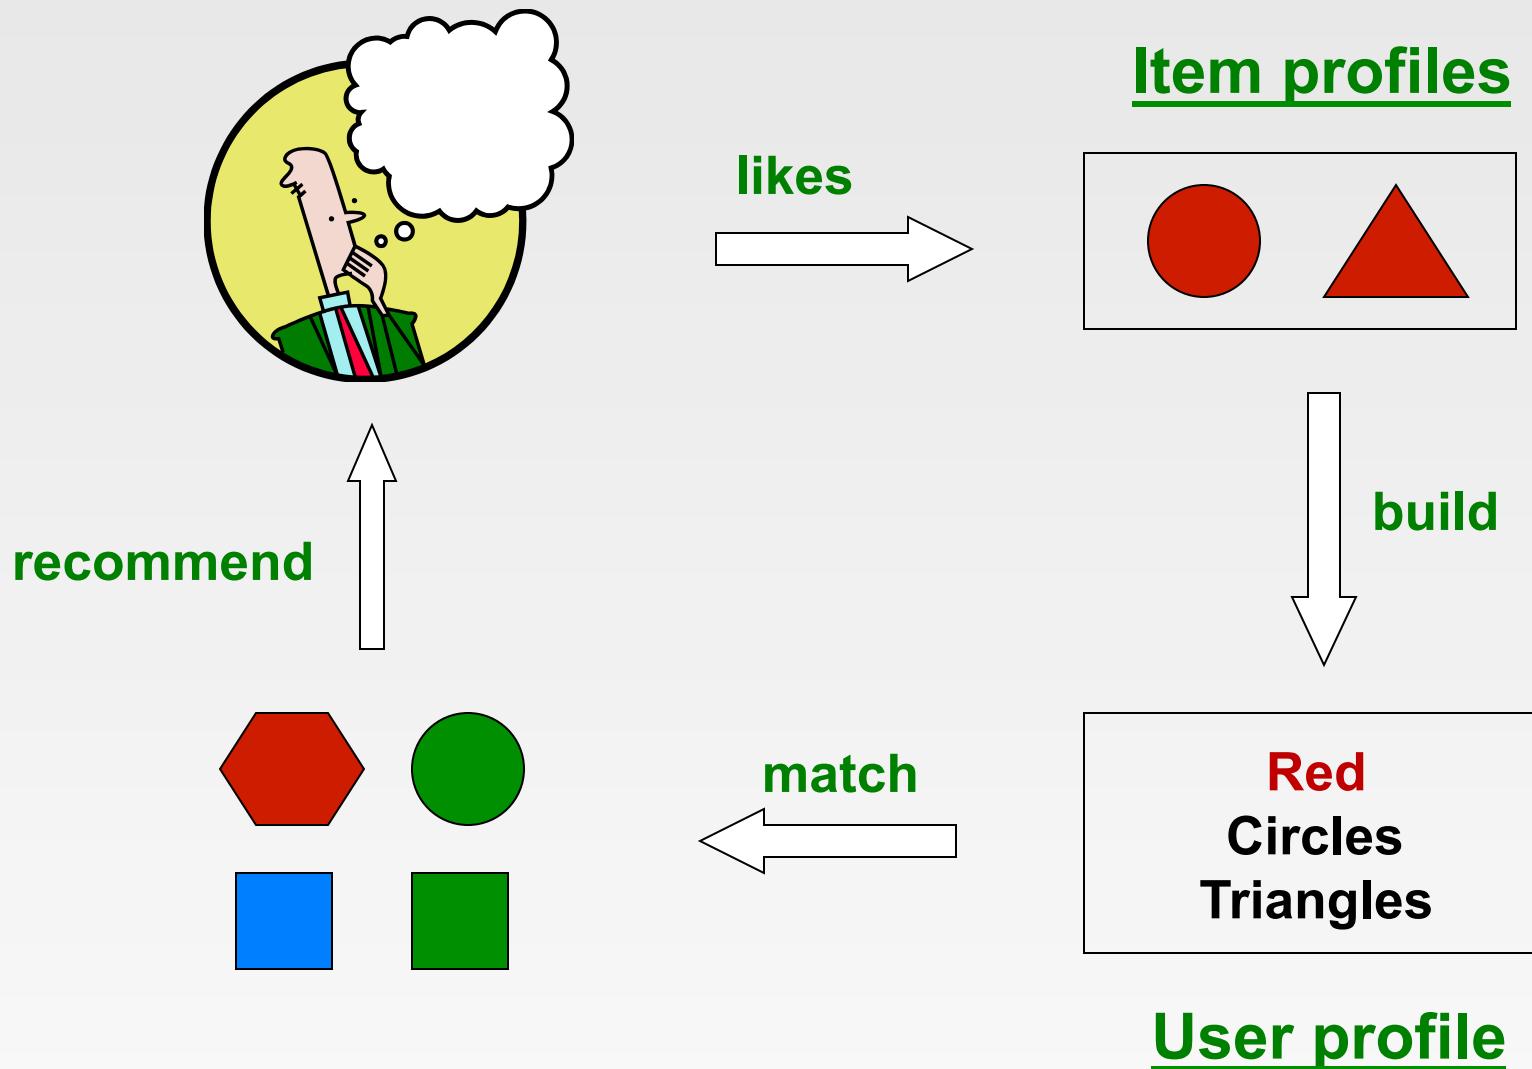
# Recommender systems: basic techniques

	Pros 	Cons 
Content-based	No community required, comparison between items possible	Content descriptions necessary, cold start for new users, no surprises
Collaborative	No knowledge-engineering effort, serendipity of results, learns market segments	Requires some form of rating feedback, cold start for new users and new items
Knowledge-based	Deterministic recommendations, assured quality, no cold-start, can resemble sales dialogue	Knowledge engineering effort to bootstrap, basically static, does not react to short-term trends

# Content-based Recommendations

- **Main idea:** Recommend items to customer  $x$  similar to previous items rated highly by  $x$
- What do we need:
  - Some information about the available items such as the genre ("content")
  - Some sort of *user profile* describing what the user likes (the preferences)
- **Example:**
  - Movie recommendations:
    - ▶ Recommend movies with same actor(s), director, genre, ...
  - Websites, blogs, news:
    - ▶ Recommend other sites with “similar” content

# Plan of Action



# Item Profiles

- For each item, create an **item profile**
- Profile is a set (vector) of features
  - **Movies:** author, title, actor, director,...
  - **Text:** Set of “important” words in document
- How to pick important features?
  - Usual heuristic from text mining is **TF-IDF**  
(Term frequency \* Inverse Doc Frequency)
    - ▶ **Term ... Feature**
    - ▶ **Document ... Item**

# Term-Frequency - Inverse Document Frequency (TF-IDF)

- Compute the overall importance of keywords

- Given a keyword  $i$  and a document  $j$

$$TF-IDF(i,j) = TF(i,j) * IDF(i)$$

- Term frequency (TF)

- Let  $freq(i,j)$  number of occurrences of keyword  $i$  in document  $j$
  - Let  $maxOthers(i,j)$  denote the highest number of occurrences of another keyword of  $j$
  - $TF(i,j) = freq(i,j) / maxOthers(i,j)$

- Inverse Document Frequency (IDF)

- $N$ : number of all recommendable documents
  - $n(i)$ : number of documents in which keyword  $i$  appears
  - $IDF(i) = \log N / n(i)$

# User Profiles and Prediction

- User profile possibilities:
  - Weighted average of rated item profiles
  - **Variation:** weight by difference from average rating for item
  - ...
  
- **Prediction heuristic:**
  - Given user profile  $x$  and item profile  $i$ , estimate
$$u(x, i) = \cos(x, i) = \frac{x \cdot i}{\|x\| \cdot \|i\|}$$

# Pros: Content-based Approach

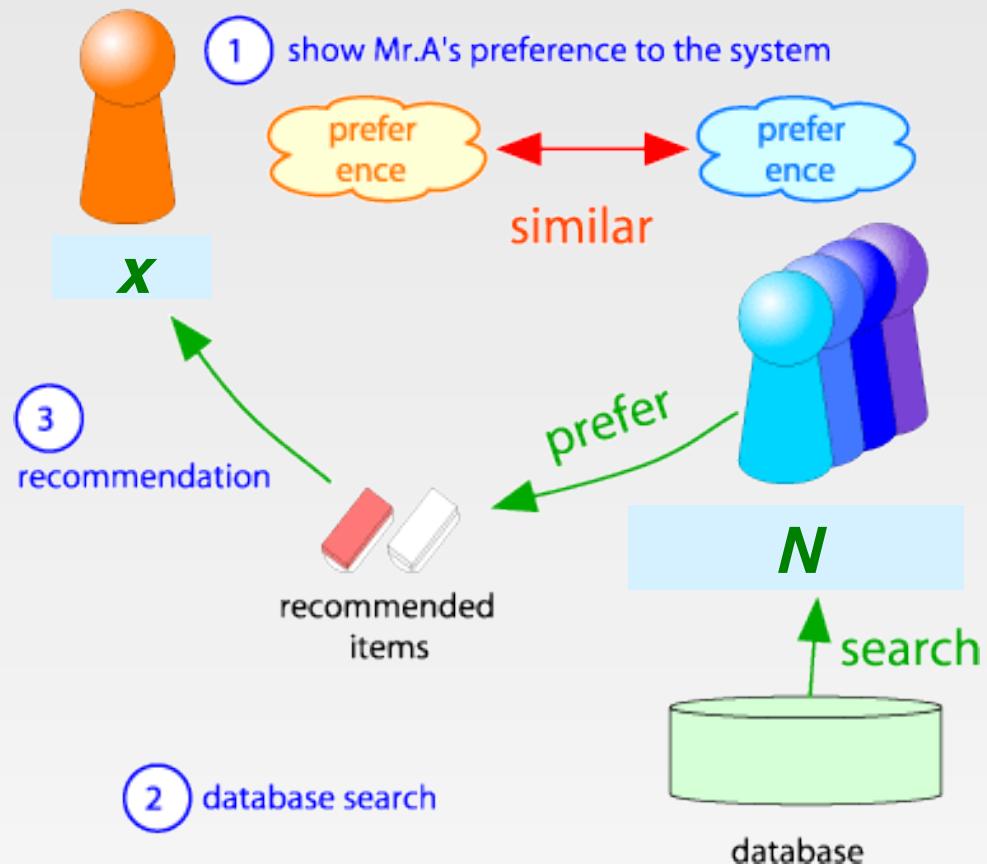
- **+: No need for data on other users**
  - No cold-start or sparsity problems
- **+: Able to recommend to users with unique tastes**
- **+: Able to recommend new & unpopular items**
  - No first-rater problem
- **+: Able to provide explanations**
  - Can provide explanations of recommended items by listing content-features that caused an item to be recommended

# Cons: Content-based Approach

- --: Finding the appropriate features is hard
  - E.g., images, movies, music
- --: Recommendations for new users
  - How to build a user profile?
- --: Overspecialization
  - Never recommends items outside user's content profile
  - People might have multiple interests
  - Unable to exploit quality judgments of other users

# Collaborative Filtering

- Consider user  $x$
- Find set  $N$  of other users whose ratings are “similar” to  $x$ 's ratings
- Estimate  $x$ 's ratings based on ratings of users in  $N$



# User-based Nearest-Neighbor Collaborative Filtering

## The basic technique:

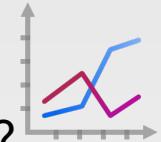
- Given an "active user" (Alice) and an item  $i$  not yet seen by Alice
- The *goal is to estimate Alice's rating for this item*, e.g., by
  - find a set of users (peers) who liked the same items as Alice in the past **and** who have rated item  $i$
  - use, e.g. the average of their ratings to predict, if Alice will like item  $i$
  - do this for all items Alice has not seen and recommend the best-rated

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

# User-based Nearest-Neighbor Collaborative Filtering (Cont')

## ■ Some first questions

- How do we measure similarity?
- How many neighbors should we consider?
- How do we generate a prediction from the neighbors' ratings?



	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

# Finding “Similar” Users

- Let  $r_x$  be the vector of user  $x$ 's ratings

$$\begin{aligned}r_x &= [* , \_, \_, *, **] \\r_y &= [* , \_, **, **, \_]\end{aligned}$$

- Jaccard similarity measure**  $\|r_x \cap r_y\| / \|r_x \cup r_y\|$

- Problem: Ignores the value of the rating

- Cosine similarity measure**

- $\text{sim}(x, y) = \cos(r_x, r_y) = r_x \cdot r_y / \|r_x\| \cdot \|r_y\|$

- Problem: Treats missing ratings as “negative”

- Pearson correlation coefficient**

- $S_{xy}$  = items rated by both users  $x$  and  $y$

$r_x, r_y$  as sets:  
 $r_x = \{1, 4, 5\}$   
 $r_y = \{1, 3, 4\}$

$r_x, r_y$  as points:  
 $r_x = \{1, 0, 0, 1, 3\}$   
 $r_y = \{1, 0, 2, 2, 0\}$

$$\text{sim}(x, y) = \frac{\sum s \in S \downarrow xy \uparrow (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum s \in S \downarrow xy \uparrow (r_{xs} - \bar{r}_x)^2} \sqrt{\sum s \in S \downarrow xy \uparrow (r_{ys} - \bar{r}_y)^2}}$$

$\bar{r}_x, \bar{r}_y$  ... avg.  
rating of  $x, y$

# Similarity Metric

Cosine similarity:

$$\text{sim}(x, y) = \frac{\sum_i r_i x_i \cdot r_i y_i}{\sqrt{\sum_i r_i^2 x_i^2} \sqrt{\sum_i r_i^2 y_i^2}}$$

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

■ Intuitively we want:  $\text{sim}(A, B) > \text{sim}(A, C)$

■ Jaccard similarity:  $1/5 < 2/4$

■ Cosine similarity:  $0.380 > 0.322$

$$\frac{4 \times 5}{\sqrt{4^2 + 5^2 + 1^2} \sqrt{5^2 + 5^2 + 4^2}} = 0.380$$

- Considers missing ratings as “negative”
- Solution: subtract the (row) mean

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

**sim A,B vs. A,C:**  
0.092 > -0.559

Notice cosine sim. is correlation when data is centered at 0

# Similarity Metric (Cont')

- A popular similarity measure in user-based CF: **Pearson correlation**

$$sim(x,y) = \frac{\sum_{s \in S} (x_s - \bar{x})(y_s - \bar{y})}{\sqrt{\sum_{s \in S} (x_s - \bar{x})^2} \sqrt{\sum_{s \in S} (y_s - \bar{y})^2}}$$

- Possible similarity values between -1 and 1;

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

sim = 0,85  
sim = 0,70  
sim = -0,79



# Rating Predictions

## From similarity metric to recommendations:

- Let  $r_x$  be the vector of user  $x$ 's ratings
- Let  $N$  be the set of  $k$  users most similar to  $x$  who have rated item  $i$
- Prediction for item  $s$  of user  $x$ :
  - $r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$
  - $r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}}$
  - Other options?
- Many other tricks possible...

Shorthand:

$$s_{xy} = sim(x, y)$$

# Memory-based and Model-based Approaches

- User-based CF is said to be "memory-based"
  - the rating matrix is directly used to find neighbors / make predictions
  - does not scale for most real-world scenarios
  - large e-commerce sites have tens of millions of customers and millions of items
- Model-based approaches
  - based on an offline pre-processing or "model-learning" phase
  - at run-time, only the learned model is used to make predictions
  - models are updated / re-trained periodically
  - large variety of techniques used
  - model-building and updating can be computationally expensive

# Item-Item Collaborative Filtering

- So far: User-user collaborative filtering
- Another view: Item-item
  - Basic idea:
    - ▶ Use the similarity between items (and not users) to make predictions
  - For item  $i$ , find other similar items
  - Estimate rating for item  $i$  based on ratings for similar items
  - Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i; x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i; x)} s_{ij}}$$

$s_{ij}$ ... similarity of items  $i$  and  $j$

$r_{xj}$ ... rating of user  $u$  on item  $j$

$N(i; x)$ ... set items rated by  $x$  similar to  $i$

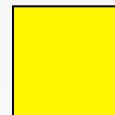
# Item-Item Collaborative Filtering

- Example:
  - Look for items that are similar to Item5
  - Take Alice's ratings for these items to predict the rating for Item5

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

# Item-Item CF ( $|N|=2$ )

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3			5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

 - unknown rating       - rating between 1 to 5

# Item-Item CF ( $|I|=2$ )

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

 - estimate rating of movie 1 by user 5

# Item-Item CF (|NI|=2)

	users												
	1	2	3	4	5	6	7	8	9	10	11	12	
	1	1	3		?	5			5		4		sim(1,m)
	2		5	4			4			2	1	3	1.00
	3	2	4		1	2		3		4	3	5	-0.18
	4		2	4		5			4		2		0.41
	5		4	3	4	2					2	5	-0.10
	6	1		3		3			2		4		-0.31
													0.59

## Neighbor selection:

Identify movies similar to movie 1, rated by user 5

Here we use Pearson correlation as similarity:

1) Subtract mean rating  $m_i$  from each movie  $i$

$$m_i = (1+3+5+5+4)/5 = 3.6$$

row 1: [-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]

2) Compute cosine similarities between rows

# Item-Item CF (|NI|=2)

	users												
	1	2	3	4	5	6	7	8	9	10	11	12	
1	1		3		?	5			5		4		$\text{sim}(1,m)$
2			5	4			4			2	1	3	1.00
3	2	4		1	2		3		4	3	5		-0.18
4		2	4		5			4			2		0.41
5			4	3	4	2					2	5	-0.10
6	6	1		3	3			2			4		-0.31
													0.59

Compute similarity weights:

$$s_{1,3}=0.41, s_{1,6}=0.59$$

# Item-Item CF (|N|=2)

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		2.6	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

Predict by taking weighted average:

$$r_{ix} = \sum_{j \in N(i,x)} s_{ij} \cdot r_{jx}$$

$$r_{1,5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

**Before:**

$$r_{xi} = \frac{\sum_{j \in N(i; x)} s_{ij} r_{xj}}{\sum_{j \in N(i; x)} s_{ij}}$$

# CF: Common Practice

- Define **similarity**  $s_{ij}$  of items  $i$  and  $j$
- Select  $k$  nearest neighbors  $N(i; x)$ 
  - Items most similar to  $i$ , that were rated by  $x$
- Estimate rating  $r_{xi}$  as the weighted average:

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i; x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i; x)} s_{ij}}$$


**baseline estimate for  $r_{xi}$**

$$b_{xi} = \mu + b_x + b_i$$

- $\mu$  = overall mean movie rating
- $b_x$  = rating deviation of user  $x$   
 $= (\text{avg. rating of user } x) - \mu$
- $b_i$  = rating deviation of movie  $i$

# Item-Item vs. User-User

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.8	
Bob		0.5		0.3
Carol	0.9		1	0.8
David			1	0.4

- In practice, it has been observed that item-item often works better than user-user
  - Why? Items are simpler, users have multiple tastes

# Pros/Cons of Collaborative Filtering

- + Works for any kind of item
  - No feature selection needed
- - Cold Start:
  - Need enough users in the system to find a match
- - Sparsity:
  - The user/ratings matrix is sparse
  - Hard to find users that have rated the same items
- - First rater:
  - Cannot recommend an item that has not been previously rated
  - New items, Esoteric items
- - Popularity bias:
  - Cannot recommend items to someone with unique taste
  - Tends to recommend popular items

# Collaborative Filtering: Complexity

- Expensive step is finding  $k$  most similar customers: **O( $|X|$ )**
- **Too expensive to do at runtime**
  - Could pre-compute
- Naïve pre-computation takes time  **$O(k \cdot |X|)$** 
  - $X \dots$  set of customers
- Ways of doing this:
  - Near-neighbor search in high dimensions (**LSH**)
  - Clustering
  - Dimensionality reduction
  - ....
- Supported by Hadoop: Apache Mahout  
<https://mahout.apache.org/users/basics/algorithms.html>

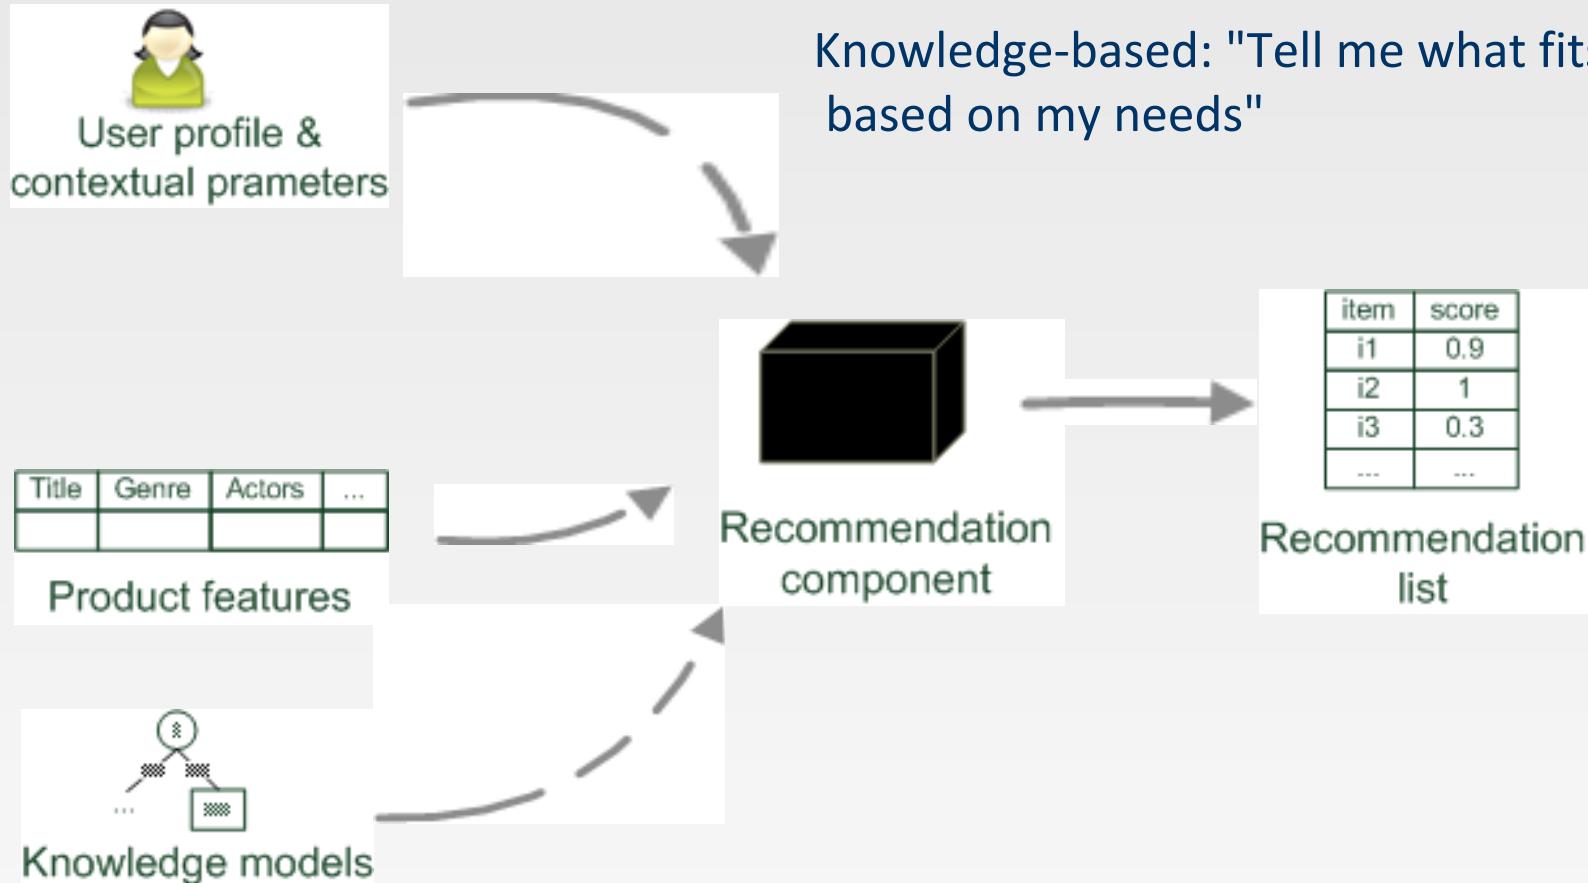
# Why do we need knowledge-based recommendation?

- Products with low number of available ratings



- Time span plays an important role
  - Five-year-old ratings for computers
  - User lifestyle or family situation changes
- Customers want to define their requirements explicitly
  - “The color of the car should be black”

# Knowledge-Based Recommendation



# Knowledge-based Recommendation

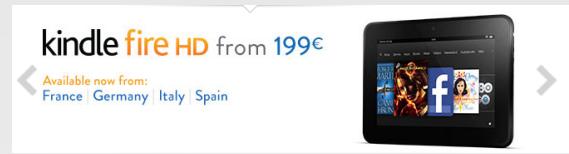
- Constraint-based
  - based on explicitly defined set of recommendation rules
  - fulfill recommendation rules
- Case-based
  - based on different types of similarity measures
  - retrieve items that are similar to specified requirements
- Both approaches are similar in their **conversational** recommendation process
  - users specify the requirements
  - systems try to identify solutions
  - if no solution can be found, users change requirements

# Hybrid Methods

- Implement two or more different recommenders and combine predictions
  - Perhaps using a linear model
- Add content-based methods to collaborative filtering
  - Item profiles for new item problem
  - Demographics to deal with new user problem

# What is a Good Recommendation in Practice?

- Total sales numbers
- Promotion of certain items
- ...
- Click-through-rates
- Interactivity on platform
- ...
- Customer return rates
- Customer satisfaction and loyalty



# Evaluation

movies					
users	1	3	4		
		3	5		5
			4	5	5
			3		
			3		
	2			2	2
					5
		2	1		1
		3			3
	1				

# Evaluation

movies					
users	1	3	4		
		3	5		5
			4	5	5
			3		
			3		
	2			?	?
					?
		2	1		?
		3		?	
	1				

Test Data Set

# Evaluating Predictions

## ■ Compare predictions with known ratings

- Root-mean-square error (RMSE)
  - ▶  $\sqrt{\frac{1}{n} \sum_{i=1}^n (r_{xi} - \hat{r}_{xi})^2}$  where  $r_{xi}$  is predicted,  $\hat{r}_{xi}$  is the true rating of  $x$  on  $i$
- Precision at top 10:
  - ▶ % of those in top 10
- Rank Correlation:
  - ▶ Spearman's *correlation* between system's and user's complete rankings

## ■ Another approach: 0/1 model

- Coverage:
  - ▶ Number of items/users for which system can make predictions
- Precision:
  - ▶ Accuracy of predictions
- Receiver operating characteristic (ROC)
  - ▶ Tradeoff curve between false positives and false negatives

# References

- Chapter 9 of Mining of Massive Datasets.
- Tutorial: Recommender Systems. IJCAI 2013.

**End of Chapter 11**